Which Side Is Wrong? Neutral Contradiction Flags Trade Misinformation for Abstention in Retrieval-Augmented Generation

Anonymous ACL submission

Abstract

Retrieval-augmented generation systems increasingly detect contradictions among retrieved passages and flag them to the generator. We ask what such a flag must say, and what its detector must get right, for it to help against misinformation. In experiments pre-registered in stages, on four benchmarks (including RAGuard's real misleading Reddit posts) and with seven generators from three model families, we vary the detector and the format of its flag. A perfectly accurate neutral flag ("these passages conflict") trades misinformation for abstention: it rarely changes which side the model picks, and its effect on accuracy varies widely across 34 settings (pooled -2.8 points; prediction interval $[-12.9, +7.4]$), so it pays only where a wrong answer costs more than no answer. An oracle that names the false side helps on every benchmark that labels one, but about as much as stating the wrong answer outright, so it is a ceiling. With a real detector the registered gains are smaller and depend on the generator, and what decides them is where the detector's errors fall: per question, a flag naming the false side raises the chance of a correct answer by 34.5 points and one marking a correct passage lowers it by 31.6. A detector's overall accuracy does not predict its value to a generator once the generator's own accuracy is accounted for, and the gain is not the detector model's answer passed along. On genuine conflicts, the same directional flag removes valid answers. We release the harness, the pre-registrations and every run.

1 Introduction

Retrieval-augmented generation (RAG) (Lewis et al., 2020) hands a language model passages that may disagree. Some disagreements are genuine: two sources report different findings, or two people share a name. Others are planted or stale misinformation (Zou et al., 2025; Wang et al., 2025b). A natural fix is to detect the contradiction and tell the generator. Recent systems pair contradiction detectors with prompts that flag conflicting passages (Wang et al., 2026), and prompting a model about conflicts can help (Cattan et al., 2025). These systems report detection F1 and end-to-end accuracy, but not how the two connect. What should the flag say, and what must the detector get right for it to help?

We fix the retrieved passages and vary, independently, the *detector* (heuristic, NLI, LLM, random placebo, oracle, and synthetic detectors of chosen quality) and the *format* that passes its output to the generator: a neutral "these passages may conflict" flag, a flag that says one side is false without saying which, a *directional* flag that names the likely misinformation, or removal of the suspect passage. After pilots on RAMDocs (Wang et al., 2025b), we pre-registered our hypotheses (H1–H9) and tested them on 300 held-out questions; seven addenda (H10–H29), each committed before its runs, extended the study to ConflictBank (Su et al., 2024) (including a variant where the misinformation outnumbers the truth), to RAGuard (Zeng et al., 2025), whose misleading side is real Reddit posts, and to seven generators from OpenAI, Anthropic and Qwen. Controls separate what a flag conveys: a forced-choice prompt with no "unknown", an answer key that names the wrong answer without pointing at any passage, the detector model's own answer given as a hint, and a crossed design in which six detectors serve six generators. WikiContradict (Hou et al., 2024), whose conflicts are genuine, is the contrast case.

Our findings:

- **A correct neutral flag trades misinformation for abstention.** It moves misinformed answers, and some correct ones, into "unknown" and barely changes accuracy on the answers that remain. Its effect on strict accuracy ranges from harmful to slightly helpful across 34 settings

(prediction interval $[-12.9, +7.4]$); when a misinformed answer is costed at four correct ones, it helps in 4 settings and hurts in 4.

- **An oracle's direction is a ceiling.** Naming the false side helps on every benchmark that labels one, but an answer key that names the wrong answer does nearly as well (H23), so the oracle rows bound what any detector can deliver.
- **With a real detector, its errors decide.** With a gpt-4o-mini detector the registered gains depend on the generator: on ConflictBank 1 of 7 generators gain significantly and 1 loses. Across 30 crossed detector–generator cells, the registered reliability gap correlates with the gain, but almost entirely through the generator's own accuracy. What matters is where the detector errs: generators follow a wrong-side flag more readily than a right one, so a detector helps a generator only if it is right where that generator is wrong and silent where it is right. The detector model's own answer as a hint does not reproduce the gain (H27), and reversing or paraphrasing the passages leaves direction unchanged (H29).
- **Direction is unsafe on genuine conflicts.** On WikiContradict the LLM directional flag, and removal, cut the share of answers that give both valid sides (H28).

For conflict-aware RAG, we recommend measuring a detector's wrong-side rate on questions the target generator already answers correctly, letting the detector abstain, passing direction only where conflicts are likely to be misinformation, and treating a neutral flag as a caution signal whose value depends on the cost of a wrong answer. We release the harness, pre-registrations and every run.[1]

2 Related Work

Knowledge conflicts. Language models can prefer memorised knowledge over contradicting context (Longpre et al., 2021), yet are readily persuaded by coherent contradicting evidence (Xie et al., 2024), to a degree that depends on the strength of their prior (Wu et al., 2024). Of the conflict types of Xu et al. (2024), we study inter-context conflicts, as in Chen et al. (2022), who find that models tend to rely on a single passage when evidence conflicts, and Wan et al. (2024), who find that models weigh a passage's relevance more than its credibility cues.

Benchmarks. RAMDocs (Wang et al., 2025b), built on AmbigDocs (Lee et al., 2024), mixes genuine ambiguity, misinformation and noise. CONFLICTS (Cattan et al., 2025) annotates search results with conflict types, each with its own desired behaviour, and WikiContradict (Hou et al., 2024) pairs genuinely contradicting Wikipedia passages. FaithEval (Ming et al., 2025), ConflictBank (Su et al., 2024) and ConfRAG (Yuan et al., 2026) cover related settings, as do QACC's human-annotated natural conflicts (Liu et al., 2025) and WhoQA's same-name conflicts (Pham et al., 2024). RAMDocs and ConflictBank label which side is misinformation; WikiContradict does not.

Telling the model about conflicts. Prompting a model to reason about the conflict type helps, and supplying the gold type helps much more (Cattan et al., 2025); on WikiContradict, models acknowledge both sides more often when told the passages may conflict (Hou et al., 2024). Ragability (Gross et al., 2026) finds that such a hint helps on contradictory contexts but hurts on *consistent* ones, where it is false. In our neutral-flag result the hint is true and the input contradictory, and accuracy still falls, because the hint does not say which side to trust. Pan et al. (2023b) test detection and "vigilant" prompting against LLM-generated misinformation, and Pan et al. (2023a) call for misinformation-aware QA. Models weigh retrieved pages' metadata and appearance (Chiang and Lee, 2024), and source-credibility signals help a fact-checker resolve conflicting evidence (Ge et al., 2025). We separate the *content* of the hint from its *accuracy*.

Telling the model which passage is wrong. Closest to our directional format, Hong et al. (2024) pass a discriminator's marks on perturbed passages to the generator in the prompt, which makes it more robust to counterfactual noise. CrAM (Deng et al., 2025) rescales attention by passage credibility, RA-RAG (Hwang et al., 2025) estimates source reliability, and Weller et al. (2024) defend open-domain QA against disinformation. TACS (Yu et al., 2024) detects untruthful context and masks it, close to our removal format. These works, with Pan et al. (2023b)'s finding that vigilant prompts give limited protection, already suggest that a reliability signal helps more than

[1] Code, the pre-registration and all run outputs are in the anonymised supplementary material and will be released publicly on acceptance.

a warning. What we add is the controlled comparison on the same pairs, its cost on conflict-free questions, and an analysis of when a real detector's signal helps a given generator. RGB (Chen et al., 2024) warns models that passages may contain errors, and they still rarely detect them; our neutral flag is the pair-specific version of that warning. Models also follow injected hints, wrong ones included (Turpin et al., 2023), and advice improves a decision-maker only when its errors fall where the decision-maker is right less often (Bansal et al., 2021), which is the pattern our crossed design measures. Whether abstaining is a failure depends on the cost of a wrong answer (Wen et al., 2025; Song et al., 2025).

Detection, resolution and poisoning. ConflictRAG (Wang et al., 2026) trains a pairwise detector with an LLM fallback and reports pair-level F1 alongside answer gains; its NLI and chain-of-thought baselines resemble two of our detectors. LLMs are imperfect context validators (Gokul et al., 2025), detect conflicts but struggle to pinpoint them (Wang et al., 2024), and tend to pick a side without explanation (Jiayang et al., 2024). They favour the more frequent side and the one matching their prior (Jin et al., 2024; Kortukov et al., 2024), prefer LLM-generated passages (Tan et al., 2024) (both ConflictBank sides are generated), are distracted by irrelevant context (Shi et al., 2023), and favour some sources (Schuster et al., 2026). TCR (Ye et al., 2026) uses a learned soft prompt, context-aware decoding (Shi et al., 2024) upweights the context, EvidentialRAG (Hossain et al., 2026) routes between answering and abstaining, and credibility-aware generation (Pan et al., 2024) fine-tunes models to weigh passage credibility, a trained analogue of our directional flags. MADAM-RAG (Wang et al., 2025b), Astute RAG (Wang et al., 2025a), Self-RAG (Asai et al., 2024), Corrective RAG (Yan et al., 2024) and InstructRAG (Wei et al., 2025) handle unreliable passages without an explicit contradiction detector; we compare against Astute RAG and TrustRAG. A few injected passages can steer a RAG system (Zou et al., 2025); RobustRAG (Xiang et al., 2024) isolates passages for certifiable robustness, and TrustRAG (Zhou et al., 2025) filters suspected poisoned ones, the analogue of our removal. We ask instead how a detector's output should be passed to the generator.

3 Framework

A question q comes with retrieved passages $P = (p_0, \ldots, p_{n-1})$. A *detector* returns flagged pairs $F \subseteq \{(i, j) : i < j\}$; a *directional* detector also returns, for each flagged pair, a *suspect* set $S_{ij} \subseteq \{i, j\}$: the side it believes is misinformation. A *format* turns (P, F, S) into the generator's prompt. We vary detectors and formats independently. Every prompt shares a skeleton: passages, question, and an instruction to end with a `Final answer(s):` line. The *default* skeleton asks for every valid answer when the passages support several (RAMDocs contains same-name ambiguity, so a correct answer may need several entities). The *minimal* skeleton drops that policy, and its answer line allows several answers without inviting them. The *single* skeleton says the question has exactly one correct answer (true of ConflictBank by construction). All three allow "unknown" when no passage answers the question; the *forced* skeleton is *single* without that option (Appendix A).

3.1 Formats

No flags (plain). Passages and question only.

Instruction. Plain plus a generic instruction to check the passages for contradictions, with no detector. v1 ends "if a disagreement reflects genuinely different valid answers, report each one"; v2 asks the model to judge which passages are reliable and to leave out answers supported only by unreliable ones.

Flag. Each flagged passage gets the header `[POTENTIAL CONFLICT with passage(s): j]`, followed by a list `Passages i and j: possible contradiction`. This is the original CAAG format, except that every detector's flags are rendered identically, so detectors differ only in *which* pairs they flag.

One-wrong flag. As Flag, but adds that exactly one passage of the pair contains misinformation, without saying which: it says that the conflict is misinformation, but not its direction.

Directional flag. As Flag, but suspect passages are headed `[LIKELY MISINFORMATION; conflicts with passage(s): j]`, and the list says which passage likely contains misinformation.

Controls. *Pick*: the neutral flag plus "decide which side is better supported and answer with

that side only". *No pairs*: each flag without its pair structure (flagged passages are marked, but not which conflicts with which). *Answer key*: plain plus "The answer X is false" for each wrong answer, with no passage marked.

Remove (suspect). The suspect passages of every flagged pair are deleted. For a pair without a suspect, the second passage is deleted, as the original CAAG REMOVE strategy does without retrieval scores.

Single call. No detector: one call that first writes `Suspect passages: ...` and then answers from the other passages, at half the calls of a detector pipeline.

Closed book. No passages; it measures what the generator knows without retrieval.

3.2 Detectors

CAAG (Contradiction-Aware Attention Gating) is a heuristic detector from an earlier, unpublished RAG system, used with its code unchanged. It scores similar pairs by negation, entity-mismatch and temporal cues, runs locally and gives no direction (Appendix B).

NLI is `nli-deberta-v3-base` (He et al., 2023), a cross-encoder scored in both directions; a pair is flagged if the larger contradiction probability exceeds 0.5. It gives no direction.

LLM is a single call that lists contradicting pairs, telling them apart from ambiguity, complementary and temporal relations, and names the side more likely to be misinformation (or "both" or "unclear"), as JSON. It is gpt-4o-mini (temperature 0) unless stated; in the crossed design gpt-4.1, gpt-6-luna, gpt-6.1-terra, Claude Haiku 5.5 and Claude Sonnet 5.5 run the same prompt.

Random flags each pair independently at CAAG's per-pair flag rate, a placebo that matches CAAG's flag volume.

Oracle flags exactly the gold contradiction pairs and names the gold misinformation side. The *inverted oracle*, an adversarial control, flags the same pairs but names the correct side.

3.3 Synthetic detectors

To separate a detector's accuracy from its identity, we build detectors from the gold labels with a chosen pair-level precision $p \in \{0.25, 0.5, 0.75, 1\}$, recall $r \in \{0.25, 0.5, 1\}$ and direction accuracy $d \in \{0.5, 0.75, 1\}$ (36 cells), each rendered as a directional flag and as removal. All cells reuse the same random draws, so they are nested and comparisons across cells are paired (Appendix E).

4 Experimental Setup

Data. **RAMDocs** (Wang et al., 2025b) has 500 questions, each with 1–12 documents (5.5 on average) typed *correct* (supports a gold answer, possibly for a different same-name entity), *misinfo* (supports a wrong answer) or *noise*. A pair is a *contradiction* if one side is misinformation and the answers differ, *ambiguity* if both are correct with different answers, *consistent* if the answers match, and *unrelated* if either is noise. A question is a *conflict* question if it has a contradiction pair, and *clean* otherwise. Some documents typed *correct* contain a wrong answer; questions whose only sign of conflict is such a document count only in "All" (Appendix D).

ConflictBank (Su et al., 2024): for a Wikidata fact, an LLM writes one evidence passage for the true object and one for a substituted object. We sample 400 questions round-robin over relations: 300 conflict questions (correct passage, misinformation passage, two distractors from other subjects) and 100 clean ones (correct passage, three distractors). Every question has one gold answer. It had no pilot, so all its results are confirmatory. A *majority* variant gives each conflict question the correct passage and ConflictBank's three wrong-object passages for that fact (including its temporal and semantic conflicts), so the misinformation outnumbers the truth three to one.

RAGuard (Zeng et al., 2025) pairs PolitiFact claims (true or false) with Reddit posts that an LLM labelled supporting, misleading or unrelated, so the misleading side is real online text. We keep 294 claims with at least one supporting and one misleading post (up to two of each, plus two unrelated posts) and 100 claims with no misleading post; the question is whether the claim is true. Its labels are LLM-assigned stances, not verified facts.

WikiContradict (Hou et al., 2024) has 253 questions, each with two genuinely contradicting Wikipedia passages; the reference answer is *both* sides. We also run 237 automatically scorable **CONFLICTS** (Cattan et al., 2025) questions (Appendix D).

Pre-registration and splits. We developed the method on a random 200-question pilot split of RAMDocs. Before running anything on the remaining 300 questions (139 conflict, 126 clean, 35 excluded), we committed a pre-registration that fixes the conditions, metrics, hypotheses and corrections. Every RAMDocs result in the main text is on this confirmatory split; pilot results appear only in the appendix. Comparisons that were not registered, including every real-detector row of Table 3, are exploratory. Seven addenda (H10–H29 and several re-analyses) were each committed before running anything they cover, so these hypotheses were formed after the original confirmatory results were known; the confirmatory core registered before any data is H1–H9. Re-analyses registered in the last addendum use data we had already seen, and are reported as such. The repository's commit history is the only timestamp. Appendix D lists every deviation.

Generators and baselines. The main generator is `gpt-4o-mini` (OpenAI et al., 2024) (temperature 0.1, at most 256 output tokens, seeds 0–2). We replicate key conditions with `gpt-4.1`, Qwen2.5-7B-Instruct (Qwen Team et al., 2024) (4-bit, local), two current reasoning models, `gpt-6-luna` (3 seeds) and `gpt-6.1-terra` (seed 0), and Anthropic's Claude Haiku 5.5 (3 seeds) and Claude Sonnet 5.5 (seed 0). The reasoning and Claude models run at their default temperature with up to 4,096 output tokens, since they reason before answering. The LLM detector is gpt-4o-mini except in the crossed design (H24), where six detectors and six generators form 30 pairs on ConflictBank. On ConflictBank the primary prompt is the single skeleton; the default and minimal skeletons are robustness checks. We rerun the oracle and LLM-detector conditions with the minimal skeleton. MADAM-RAG (Wang et al., 2025b) is ported unchanged (three debate rounds). Astute RAG (Wang et al., 2025a) and TrustRAG (Zhou et al., 2025) use their published prompts and defaults; since both return one answer, we also run *multi* variants that ask for every valid answer (Appendix D). We also run the LLM detector on Qwen2.5-7B with gpt-4o-mini generating. A detector pipeline counts as one detector call plus one generator call. Two controls test whether a crossed cell's gain is only a stronger model's answer (H27): the generator given the detector model's own plain answer as a hint, and the detector model used as the generator. An artefact probe (H29) reruns two detectors on ConflictBank with the passage order reversed and with every passage paraphrased by one model.

Metrics. On RAMDocs an answer is correct if it includes every gold answer and no wrong answer (Wang et al., 2025b). Our primary metric reads "includes" literally: after normalisation, each gold or wrong answer must appear as a whole-word span of the extracted final answer. A bidirectional variant, which also counts a predicted answer contained in a gold one, is the robustness metric. We chose the one-way metric before the confirmatory run because it is closer to MADAM-RAG's published score (Appendix D), and we compare conditions only within our own harness. On WikiContradict an answer is correct if it contains both sides (bidirectional matching); on ConflictBank and RAGuard we use the RAMDocs metrics. We also report the share of answers that contain the misinformation, the share that abstain ("unknown"), and accuracy among answered questions. An LLM judge (gpt-6.1-terra) agrees with the primary metric on 92.6% of 162 stratified conflict-question outputs ($\kappa = 0.85$); the metric is slightly conservative, missing some paraphrases. A 200-output sheet for a human audit is released, but it has not been labelled.

Statistics. We average each question's score over seeds and report paired differences from the plain format with bootstrap 95% CIs that resample questions (10,000 resamples) and two-sided sign-flip permutation tests (Efron and Tibshirani, 1993; Dror et al., 2018) (10,000 permutations, extended to 200,000 when none reaches the observed statistic, so that a p-value can survive a large Holm family). CIs are unadjusted; significance is decided by the adjusted p below, so an interval can exclude zero for a difference that is not significant. A hypothesis is *robustly supported* if it also holds under the robustness metric. One significance scheme governs the paper: a registered test takes its family's Holm-adjusted p (Holm, 1979), and every other comparison we report (636 of them) forms one exploratory Holm family. Only results significant under this scheme are called significant or set in bold. A null difference is not read as equivalence without a TOST (Lakens, 2017) at ± 3 points. Selective accuracy is a ratio (correct over answered, pooled over questions and seeds), tested by swapping conditions within questions. Pooled

estimates use a bootstrap that resamples questions jointly across the settings that share them. Appendix D gives the minimum detectable effect.

Cost. Token counts come from the API's usage reports. All experiments, pilots included, cost $185.14 in API fees (prices checked on 2026-10-27; Appendix D); Qwen ran locally.

5 Results

Tables 2 and 3 give every condition, and the hypothesis tables every registered test (Appendix B).

5.1 A neutral flag trades misinformation for abstention

Flagging exactly the gold contradictions in the neutral format lowers gpt-4o-mini's RAMDocs conflict-question accuracy by 6.7 points [−10.8, −3.1] from 22.1% (H2), and gpt-4.1's by 8.6 (H9). On ConflictBank, with a prompt stating that the question has one answer, it costs gpt-4o-mini 18.7 points, gpt-4.1 8.0 and gpt-6-luna 9.4 (H15, H19), but it does not significantly change accuracy for gpt-6.1-terra, Qwen2.5-7B, either Claude model (H25) or any generator on RAGuard (H26). Pooled over the 34 settings in which we ran it (Table 30; registered as exploratory), the effect is −2.8 points, but settings differ more than sampling explains: a new setting's effect would fall in [−12.9, +7.4]. A meta-regression places the harm in OpenAI models on generated data with "unknown" allowed (−5.7 points); Claude (+6.3), a prompt without "unknown" (+5.4) and natural data (+2.1) each shift the effect up.

The flag changes whether the model answers, not which side it picks. With the one-answer prompt, of gpt-4o-mini's answers that contained the misinformation without a flag, 45% abstain with the neutral flag and only 7% become correct; 37% of its correct answers also turn into abstentions (Table 31). Accuracy on the answers it still gives moves by −0.3 to +3.9 points across the eight ConflictBank settings of H21, significantly in 1 (Claude Haiku, forced prompt); equivalence within ±3 points is established in 1. When "unknown" is not allowed (H20), the neutral flag has no significant effect for gpt-4o-mini (−0.2), gpt-6-luna (+1.6) or gpt-6.1-terra (−0.3), though equivalence is not established, and a small one for Claude Haiku (+3.9). Under the multi-answer prompt it makes the model list both sides (Figure 3); on WikiContradict, where both are valid, that raises accuracy (H8: +6.5).

Whether that trade helps depends on what a wrong answer costs. Strict accuracy counts an abstention and a misinformed answer alike. If a misinformed answer instead costs c correct ones (abstaining costs nothing), the neutral flag significantly lowers utility in 8 of 28 settings at $c = 0$ and raises it in 1; at $c = 4$ it raises it in 4 and lowers it in 4 (Table 31); for gpt-4o-mini on ConflictBank it moves from −18.7 to +34.7.

5.2 Naming the false side: what an oracle can deliver

With the same oracle pairs, the directional flag beats the neutral and one-wrong flags in every registered comparison (H3, H9–H12, H15, H16, H19, H25, H26). On ConflictBank it gains between 17.0 and 36.4 points for the OpenAI generators and 16.3 to 23.3 for Claude; it gains under the forced prompt (H20: +14.7 to +26.4), when the misinformation outnumbers the truth three to one (+32.7 to +66.3), and on RAGuard's real posts (H26: +6.8 to +20.4). On RAMDocs it gains 13.4 points with gpt-4o-mini and up to 39.6 with gpt-6.1-terra, but not significantly for Qwen2.5-7B (Table 19).

These rows are a ceiling, not a detector result. Telling the model "the answer X is false", with no passage marked, recovers most of the gain: no registered comparison differs significantly (H23), the directional flag's advantage ranges from −5.0 to +5.3 points, and within ±3 points equivalence holds in 2 of 9 (TOST); for gpt-4o-mini on ConflictBank the flag is ahead (90% CI [+2.2, +8.6]). An inverted oracle that names the correct side costs 7.0 points on RAMDocs.

Saying only that one side is false helps when the model must answer. Where it may abstain, the one-wrong flag lowers gpt-4o-mini's ConflictBank accuracy by 35.7 points (H16); under the forced prompt it raises every generator's (H20: +5.1 to +10.3), and on RAGuard it helps Claude Haiku (+5.9) but hurts gpt-6-luna (−7.0). Removing the pair structure has no consistent effect (H22).

5.3 With a real detector, where its errors fall decides

Registered results. With the gpt-4o-mini LLM detector, the directional flag significantly raises accuracy over all ConflictBank questions for 1 of

	No flags				Oracle, neutral flag				Oracle, directional flag			
Generator	Correct	Misinfo.	Abstain	Sel.	Correct	Misinfo.	Abstain	Sel.	Correct	Misinfo.	Abstain	Sel.
ConflictBank, one-answer prompt												
gpt-4o-mini	51.1	31.0	6.7	54.8	32.4	**17.7**	**41.0**	55.0	81.2	**4.2**	**2.1**	**83.0**
gpt-6-luna	53.6	24.3	11.9	60.8	44.1	**17.4**	**28.9**	62.0	90.0	**0.0**	**0.8**	**90.7**
Claude Haiku 5.5	66.8	22.0	1.2	67.6	67.4	22.1	0.9	68.0	90.1	**0.2**	1.1	**91.1**
gpt-6.1-terra	73.0	18.0	1.7	74.2	70.7	15.7	4.7	74.1	90.0	**0.7**	0.3	**90.3**
Claude Sonnet 5.5	75.7	15.7	1.0	76.4	78.7	12.0	1.7	80.0	92.0	1.0	1.7	93.6
gpt-4.1	64.7	27.3	1.7	65.8	56.7	21.7	14.0	65.9	91.3	0.3	1.7	92.9
Qwen2.5-7B	50.3	35.3	5.0	53.0	47.0	28.7	15.0	55.3	75.7	13.0	1.3	76.7
ConflictBank, forced choice												
gpt-4o-mini	55.4	31.6	0.1	55.5	55.2	31.9	0.4	55.5	79.7	**6.7**	0.0	**79.7**
gpt-6-luna	64.6	25.8	0.0	64.6	66.1	22.0	0.0	66.1	91.0	**0.1**	0.0	**91.0**
Claude Haiku 5.5	69.3	20.9	0.0	69.3	73.2	**16.8**	0.0	**73.2**	90.4	**0.0**	0.0	**90.4**
gpt-6.1-terra	77.7	14.7	0.0	77.7	77.3	13.3	0.0	77.3	92.3	**0.3**	0.0	**92.3**
RAGuard (real posts)												
gpt-4o-mini	36.7	31.2	32.1	54.1	39.9	27.8	32.3	59.0	46.6	**18.8**	34.6	**71.2**
gpt-6-luna	23.2	20.0	56.8	53.8	23.2	21.1	55.7	52.4	32.1	**9.8**	58.2	**76.7**
Claude Haiku 5.5	24.5	17.9	57.6	57.8	24.8	20.1	55.1	55.3	37.2	**11.0**	51.8	**77.2**
gpt-6.1-terra	24.8	14.3	60.9	63.5	27.2	16.7	56.1	62.0	31.6	**8.8**	59.5	78.2
Claude Sonnet 5.5	27.2	12.6	60.2	68.4	31.6	14.6	53.7	68.4	47.6	10.9	**41.5**	81.4
ConflictBank, misinformation majority												
gpt-4o-mini	21.6	64.1	3.2	22.3	22.7	**40.7**	**26.8**	**31.0**	54.2	**21.3**	**11.7**	**61.4**
gpt-6-luna	29.8	58.8	4.2	31.1	36.3	**43.6**	**11.7**	**41.1**	90.2	**0.2**	1.7	**91.8**
Claude Haiku 5.5	31.0	58.4	1.1	31.3	35.3	55.1	0.9	35.7	88.3	**1.0**	1.8	**89.9**
gpt-6.1-terra	26.0	65.3	1.0	26.3	22.3	71.0	0.7	22.5	92.3	**0.0**	1.3	**93.6**

Table 1: Risk and coverage on conflict questions (%): strictly correct answers, answers that contain the misinformation, abstentions ("unknown"), and selective accuracy (correct answers over non-abstained answers, pooled over questions and seeds, the estimand of §15.4.5). **Bold**: the change from no flags is significant (H21 for ConflictBank; the paper's single scheme elsewhere). Where abstention is allowed, the neutral flag moves answers from misinformation, and from correct answers, into abstention, and leaves selective accuracy nearly unchanged; under forced choice it has no abstention to trade.

7 generators (gpt-6-luna, +6.1) and significantly lowers it for 1 (gpt-6.1-terra, −8.8; H15, H19, H25). It helps all 4 generators tested under majority misinformation (+5.7 to +9.0), and on RAGuard it helps the two Claude models (H26: +5.0, +7.1) but none of the OpenAI ones.

The registered reliability test does not isolate the detector. H24 crossed six detectors with six generators in 30 cells. As registered, the gain rises with the detector's direction accuracy minus the generator's no-flag accuracy ($\rho = +0.64$, p<0.001), including on the 25 cells that did not exist when H24 was registered ($\rho = +0.79$). But the generator's accuracy G enters both the predictor and the gain, and most of the correlation comes from it. Regressing the flagged accuracy on direction accuracy D and G gives D a coefficient of 0.09 against 0.46 for G; with generator fixed effects, permuting D across the six detectors gives $p = 0.29$ (720 permutations). Leaving out one detector at a time, the registered gap predicts the held-out cells' gains worse than their mean (R^2 =-0.21), and worse than G alone (0.17; Table 29). A detector's overall direction accuracy says little about whether it will help.

Where its errors fall does. Per question, with detector, generator and question effects, a flag that names the false side raises the chance of a correct answer by 34.5 points [+32.8, +36.1], and one that marks a correct passage lowers it by 31.6 [−33.9, −29.4]. Generators follow a right flag they need on 57–75% of questions and a wrong-side flag on 76–88%: they act on a wrong flag more readily than on a right one. A cell's gain is therefore set by how often the detector is right where the generator is wrong and wrong where the generator is right. The exploratory score built from these two counts predicts held-out detectors' cells ($R^2 = 0.62$; calibration slope 0.64 [0.54, 0.74]; Figure 1, left). The gpt-4o-mini detector flags almost every question (97.7%) and marks a correct passage on 30.3% of them: it helps the weaker generators and lowers the accuracy of the four most accurate, significantly only for gpt-6.1-

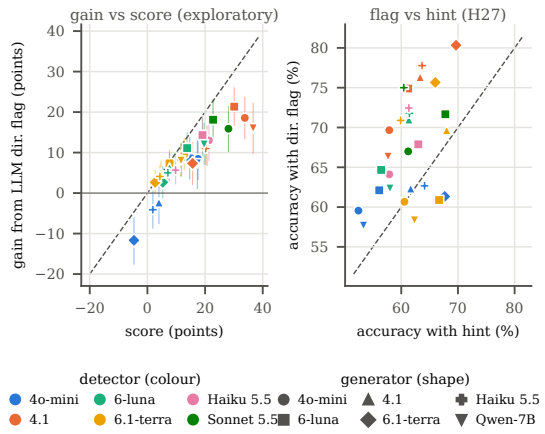


Figure 1: The 30 detector–generator cells (ConflictBank, one-answer prompt, 300 conflict questions). Left: each cell's gain from the LLM directional flag against the exploratory score (right direction where the generator errs minus a correct passage marked where it is right); dashed: $y = x$. Right: the generator's accuracy with the detector's directional flag against its accuracy given the detector model's own answer as a hint (H27); dashed: $y = x$. Bars: 95% CIs. Table 28 lists the cells.

terra (H19). gpt-6.1-terra flags 37.7% and marks a correct passage on 4.0%; it lowers no generator's accuracy.

It is not just a stronger model's answer (H27). If a cell's gain were the detector model's better answer passed along, giving the generator that answer as a hint would do as well. It does not: the directional flag beats the hint significantly in 14 of 30 cells and loses to it in 0 (Figure 1, right). The hint itself helps in 11 cells and hurts in 6, mostly the more accurate generators, which lose accuracy even when the hinted answer is better than their own. Using the detector model as the generator instead does significantly worse than the cell in 16 cells and better in 4, each a much stronger detector model (gpt-6.1-terra or Claude Sonnet) serving a weaker generator.

Not a surface artefact (H29). ConflictBank's two sides are written by an LLM from different prompts, so a detector could key on style or position. Reversing the passage order or paraphrasing every passage with one model leaves gpt-4o-mini's direction accuracy at 65.7% and 68.0% (from 67.0%) and gpt-6-luna's at 44.0% and 44.3% (from 45.7%); no change is significant, and neither is the change in the generator's gain (H29c).

Genuine conflicts (H28). Where both sides are valid, direction does damage. On WikiContradict the LLM directional flag lowers the share of answers giving both sides by 6.5 points for gpt-4o-mini, 29.0 for gpt-6-luna and 27.7 for Claude Haiku (from 14.0, 33.2 and 31.2%), and removal by more; the oracle neutral flag raises it. On the 99 RAMDocs questions whose conflicts are same-name ambiguity, removal costs every generator 21.2 to 29.3 points and the directional flag up to 14.8 (gpt-6-luna). The detector calls a valid alternative misinformation, and the generator drops it.

Baselines. Astute RAG (Wang et al., 2025a) and TrustRAG (Zhou et al., 2025) lose legitimate same-name answers on RAMDocs (H14), and a single call that names suspect passages before answering does not differ significantly from the detector pipeline (H18; Table 27). The synthetic grid (H4, H6) is in Appendix E.

6 Discussion

A neutral flag is a caution signal. It says that passages disagree but not which to trust, so the model falls back on its answer policy and often abstains. That is harm as accuracy and a trade as risk, and abstaining is right where a wrong answer is costly (Wen et al., 2025); it should be measured with a cost-sensitive score.

What a detector should optimise. A directional flag replaces part of the generator's judgement with the detector's, and generators follow a flag that marks a correct passage more readily than one they need. A detector's value to a generator is therefore set by two counts on held-out conflicts: how often it names the false side where the generator is wrong, and how often it marks a correct passage where the generator is right. Detection F1 and direction accuracy mix both. This is the complementarity condition of human–AI teams (Bansal et al., 2021); the detectors that served every generator well stayed silent when unsure.

Recommendations. Measure a detector's wrong-side rate where the target generator is already right, and let it abstain; a much stronger detector model may do better answering alone. Pass direction, or remove passages, only where a conflict is likely to be misinformation.

Limitations

Models. We test seven generators from three families: four OpenAI models (gpt-4o-mini, gpt-4.1 and two current reasoning models), Claude Haiku 5.5 and Sonnet 5.5, and Qwen2.5-7B at 4-bit precision. We could not run Gemini, and no large open model. Qwen2.5-7B, gpt-6.1-terra and Claude Sonnet ran with one seed, as did gpt-4.1 outside RAMDocs, and the reasoning and Claude models only at their default temperature, so their sampling is not deterministic, and Claude ran through Anthropic's OpenAI-compatible endpoint. The crossed reliability test (H24) covers 30 detector–generator cells on one benchmark, with six detectors, so any test of a detector-level property has little power (an exact permutation has 720 arrangements); the registered gap turned out to be dominated by the generator's own accuracy, and the score that does predict was chosen after seeing the cells and is validated only by leaving out one detector at a time on the same data. The LLM detector is a single prompt. The published baselines are reimplemented (Astute RAG has no official code), and we do not compare against a trained detector such as ConflictRAG's (Wang et al., 2026).

Data and power. The confirmatory evidence comes from two benchmarks: RAMDocs (300 questions, 139 of them conflict questions) and ConflictBank (400 sampled questions, 300 with conflicts). In both, the misinformation is synthetic, made by editing or generating text around a substituted answer; ConflictBank's passages are LLM-written on both sides. RAGuard's misleading side is real, but its stance labels are assigned by an LLM, and with its posts in context most generators abstain on most claims, which limits what flags can show there. RAMDocs' document labels are imperfect. We exclude 35 holdout questions whose status depends on documents we could not label without human annotation, and we check that the conclusions survive treating those documents as misinformation. Our comparisons detect effects of about 5.8 points with 80% power, so smaller effects, including small harms on clean questions, may be missed. WikiContradict has genuine conflicts but no misinformation side. CONFLICTS shows no significant effect for any condition; scored only on items with a reference answer, it can neither confirm nor refute our findings. The majority variant tests one case where the false side is more numerous, but its extra wrong-side passages are generated for the same substitution, which a real attacker would vary. The current reasoning models distrust ConflictBank's generated passages when asked to judge them (Appendix C), which suggests that generated benchmarks may not measure how such models treat real documents. Reversing and paraphrasing the passages leave the detectors' direction unchanged (H29), which rules out position and surface style as its source but not deeper regularities of generated text.

Metric. RAMDocs ships no scorer. Our primary metric reconstructs the published definition and scores our MADAM-RAG port close to the published number on the pilot split, but the published subset is unknown, so this is not a reproduction. The metric is string matching: it misses paraphrases, and partial answers interact with "no wrong answer". An LLM judge agrees with it on 92.6% of sampled outputs, but no human labels check either; the 200-output audit sheet we release is unlabelled. We report a more lenient variant throughout and mark any conclusion that depends on the choice.

Direction and the grid. Oracle direction is not distinguishable from an explicit answer key (H23), so the oracle rows are ceilings, not findings about detectors. In the synthetic grid the gain rises with direction accuracy, but the grid cannot separate adjacent levels. The real LLM detector names the right side 88.3% of the time on RAMDocs and 68.8% on ConflictBank. Its directional flag's gain over all ConflictBank questions with gpt-4o-mini is not significant after correction (H15), and with stronger generators it hurts. Synthetic false flags are spread uniformly and their suspect side is random, whereas real detectors concentrate their errors. The grid predicts real detectors only roughly, and three detectors cannot show whether document- or pair-level summaries predict better.

Pre-registration. The original hypotheses (H1–H9) were registered before the confirmatory run. H10–H29 were registered in seven addenda, after the earlier results were known though before their own runs, and the only timestamp is the repository's commit history, not a third-party registry. Several analyses were chosen after seeing data and are labelled as such: the TOST margin, the single- and multi-answer split, the document-level placement axes, the post-hoc comparison of detec-

tor and generator reliability that H24 later tested, H24's mechanism score, and the refinement of permutation p-values at their floor.

Scope. We study flags passed to the generator as text. Learned interfaces such as soft prompts or credibility fine-tuning (Ye et al., 2026; Pan et al., 2024), and retrieval-side filtering, may follow different trade-offs. We test one adversarial case, a detector that always names the wrong side, but not an attacker who controls which pairs are flagged.

Ethics Statement

This work uses public benchmarks and commercial model APIs, and involves no human subjects or personal data. Its main risk is dual use. Our results show that a false "likely misinformation" label, or the removal of a passage, can change a RAG system's answers: an inverted oracle that names the correct side as misinformation costs 7.0 points on RAMDocs and 8.6 on ConflictBank as a flag, and 21.8 on RAMDocs when the passage it names is removed. An adversary who can trigger a deployed detector could exploit this (Zou et al., 2025). We therefore report harm on conflict-free questions alongside every gain, and we argue that systems should measure what a detector does at the document level before acting on its flags.

References

Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, and Hannaneh Hajishirzi. 2024. Self-RAG: Learning to retrieve, generate, and critique through self-reflection. In *The Twelfth International Conference on Learning Representations (ICLR)*.

Gagan Bansal, Tongshuang Wu, Joyce Zhou, Raymond Fok, Besmira Nushi, Ece Kamar, Marco Tulio Ribeiro, and Daniel S. Weld. 2021. Does the whole exceed its parts? the effect of AI explanations on complementary team performance. In *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems*.

Dallas Card, Peter Henderson, Urvashi Khandelwal, Robin Jia, Kyle Mahowald, and Dan Jurafsky. 2020. With little power comes great responsibility. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 9263–9274, Online. Association for Computational Linguistics.

Arie Cattan, Alon Jacovi, Ori Ram, Jonathan Herzig, Roee Aharoni, Sasha Goldshtein, Eran Ofek, Idan Szpektor, and Avi Caciularu. 2025. DRAGged into conflicts: Detecting and addressing conflicting sources in search-augmented LLMs. *Preprint*, arXiv:2506.08500.

Hung-Ting Chen, Michael Zhang, and Eunsol Choi. 2022. Rich knowledge sources bring complex knowledge conflicts: Recalibrating models to reflect conflicting evidence. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 2292–2307, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.

Jiawei Chen, Hongyu Lin, Xianpei Han, and Le Sun. 2024. Benchmarking large language models in retrieval-augmented generation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 17754–17762.

Cheng-Han Chiang and Hung-yi Lee. 2024. Do metadata and appearance of the retrieved webpages affect LLM's reasoning in retrieval-augmented generation? In *Proceedings of the 7th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP*, pages 389–406, Miami, Florida, US. Association for Computational Linguistics.

Boyi Deng, Wenjie Wang, Fengbin Zhu, Qifan Wang, and Fuli Feng. 2025. CrAM: Credibility-aware attention modification in LLMs for combating misinformation in RAG. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 23760–23768.

Rotem Dror, Gili Baumer, Segev Shlomov, and Roi Reichart. 2018. The hitchhiker's guide to testing statistical significance in natural language processing. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1383–1392, Melbourne, Australia. Association for Computational Linguistics.

Bradley Efron and Robert J. Tibshirani. 1993. *An Introduction to the Bootstrap*, volume 57 of *Monographs on Statistics and Applied Probability*. Chapman & Hall, New York.

Ziyu Ge, Yuhao Wu, Daniel Wai Kit Chin, Roy Ka-Wei Lee, and Rui Cao. 2025. Resolving conflicting evidence in automated fact-checking: A study on retrieval-augmented LLMs. In *Proceedings of the Thirty-Fourth International Joint Conference on Artificial Intelligence, IJCAI-25*, pages 9656–9664. International Joint Conferences on Artificial Intelligence Organization.

Vignesh Gokul, Srikanth Tenneti, and Alwarappan Nakkiran. 2025. Contradiction detection in RAG systems: Evaluating LLMs as context validators for improved information consistency. *Preprint*, arXiv:2504.00180.

Stephanie Gross, Johann Petrak, and Brigitte Krenn. 2026. Ragability benchmark: A dataset and library to test LLMs on inter-context conflicts. In *Proceedings of the Fifteenth Language Resources and*

Evaluation Conference, pages 2323–2333, Palma de Mallorca, Spain. ELRA Language Resource Association.

Suchin Gururangan, Swabha Swayamdipta, Omer Levy, Roy Schwartz, Samuel R. Bowman, and Noah A. Smith. 2018. Annotation artifacts in natural language inference data. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 2 (Short Papers)*, pages 107–112, New Orleans, Louisiana. Association for Computational Linguistics.

Pengcheng He, Jianfeng Gao, and Weizhu Chen. 2023. DeBERTaV3: Improving DeBERTa using ELECTRA-style pre-training with gradient-disentangled embedding sharing. In *The Eleventh International Conference on Learning Representations (ICLR)*.

Sture Holm. 1979. A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70.

Giwon Hong, Jeonghwan Kim, Junmo Kang, Sung-Hyon Myaeng, and Joyce Jiyoung Whang. 2024. Why so gullible? enhancing the robustness of retrieval-augmented models against counterfactual noise. In *Findings of the Association for Computational Linguistics: NAACL 2024*, pages 2474–2495, Mexico City, Mexico. Association for Computational Linguistics.

S M Asif Hossain, Ruksat Khan Shayoni, and M. F. Mridha. 2026. EvidentialRAG: Quantifying and mitigating information conflict in multi-source retrieval-augmented generation via evidential deep learning. *Preprint*, arXiv:2607.10491.

Yufang Hou, Alessandra Pascale, Javier Carnerero-Cano, Tigran Tchrakian, Radu Marinescu, Elizabeth Daly, Inkit Padhi, and Prasanna Sattigeri. 2024. WikiContradict: A benchmark for evaluating LLMs on real-world knowledge conflicts from Wikipedia. In *Advances in Neural Information Processing Systems (Datasets and Benchmarks Track)*, volume 37, pages 109701–109747.

Jeongyeon Hwang, Junyoung Park, Hyejin Park, Dongwoo Kim, Sangdon Park, and Jungseul Ok. 2025. Retrieval-augmented generation with estimation of source reliability. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 34279–34303, Suzhou, China. Association for Computational Linguistics.

Cheng Jiayang, Chunkit Chan, Qianqian Zhuang, Lin Qiu, Tianhang Zhang, Tengxiao Liu, Yangqiu Song, Yue Zhang, Pengfei Liu, and Zheng Zhang. 2024. ECON: On the detection and resolution of evidence conflicts. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 7816–7844, Miami, Florida, USA. Association for Computational Linguistics.

Zhuoran Jin, Pengfei Cao, Yubo Chen, Kang Liu, Xiaojian Jiang, Jiexin Xu, Li Qiuxia, and Jun Zhao. 2024. Tug-of-war between knowledge: Exploring and resolving knowledge conflicts in retrieval-augmented language models. In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, pages 16867–16878, Torino, Italia. ELRA and ICCL.

Tomáš Kočiský, Jonathan Schwarz, Phil Blunsom, Chris Dyer, Karl Moritz Hermann, Gábor Melis, and Edward Grefenstette. 2018. The NarrativeQA reading comprehension challenge. *Transactions of the Association for Computational Linguistics*, 6:317–328.

Evgenii Kortukov, Alexander Rubinstein, Elisa Nguyen, and Seong Joon Oh. 2024. Studying large language model behaviors under context-memory conflicts with real documents. In *First Conference on Language Modeling*.

Daniël Lakens. 2017. Equivalence tests: A practical primer for t tests, correlations, and meta-analyses. *Social Psychological and Personality Science*, 8(4):355–362.

Yoonsang Lee, Xi Ye, and Eunsol Choi. 2024. AmbigDocs: Reasoning across documents on different entities under the same name. In *First Conference on Language Modeling (COLM)*.

Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems*, volume 33, pages 9459–9474. Curran Associates, Inc.

Siyi Liu et al. 2025. Open domain question answering with conflicting contexts. In *Findings of the Association for Computational Linguistics: NAACL 2025*, pages 1838–1854.

Shayne Longpre, Kartik Perisetla, Anthony Chen, Nikhil Ramesh, Chris DuBois, and Sameer Singh. 2021. Entity-based knowledge conflicts in question answering. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 7052–7063, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.

Yifei Ming, Senthil Purushwalkam, Shrey Pandit, Zixuan Ke, Xuan-Phi Nguyen, Caiming Xiong, and Shafiq Joty. 2025. FaithEval: Can your language model stay faithful to context, even if "the moon is made of marshmallows". In *The Thirteenth International Conference on Learning Representations (ICLR)*.

OpenAI, Aaron Hurst, Adam Lerer, Adam P. Goucher, Adam Perelman, Aditya Ramesh, et al. 2024. GPT-4o system card. *Preprint*, arXiv:2410.21276.

Liangming Pan, Wenhu Chen, Min-Yen Kan, and William Yang Wang. 2023a. Attacking open-domain question answering by injecting misinformation. In *Proceedings of the 13th International Joint Conference on Natural Language Processing and the 3rd Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 525–539, Nusa Dua, Bali. Association for Computational Linguistics.

Ruotong Pan, Boxi Cao, Hongyu Lin, Xianpei Han, Jia Zheng, Sirui Wang, Xunliang Cai, and Le Sun. 2024. Not all contexts are equal: Teaching LLMs credibility-aware generation. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 19844–19863, Miami, Florida, USA. Association for Computational Linguistics.

Yikang Pan, Liangming Pan, Wenhu Chen, Preslav Nakov, Min-Yen Kan, and William Wang. 2023b. On the risk of misinformation pollution with large language models. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 1389–1403, Singapore. Association for Computational Linguistics.

Quang Hieu Pham, Hoang Ngo, Anh Tuan Luu, and Dat Quoc Nguyen. 2024. Who's who: Large language models meet knowledge conflicts in practice. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 10142–10151.

Qwen Team, An Yang, Baosong Yang, et al. 2024. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*.

Nils Reimers and Iryna Gurevych. 2019. Sentence-BERT: Sentence embeddings using Siamese BERT-networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 3982–3992, Hong Kong, China. Association for Computational Linguistics.

Jakob Schuster, Vagrant Gautam, and Katja Markert. 2026. Whose facts win? LLM source preferences under knowledge conflicts. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 29430–29459, San Diego, California, United States. Association for Computational Linguistics.

Freda Shi, Xinyun Chen, Kanishka Misra, Nathan Scales, David Dohan, Ed H. Chi, Nathanael Schärli, and Denny Zhou. 2023. Large language models can be easily distracted by irrelevant context. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 31210–31227. PMLR.

Weijia Shi, Xiaochuang Han, Mike Lewis, Yulia Tsvetkov, Luke Zettlemoyer, and Wen-tau Yih. 2024. Trusting your evidence: Hallucinate less with context-aware decoding. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 2: Short Papers)*, pages 783–791, Mexico City, Mexico. Association for Computational Linguistics.

Maojia Song, Shang Hong Sim, Rishabh Bhardwaj, Hai Leong Chieu, Navonil Majumder, and Soujanya Poria. 2025. Measuring and enhancing trustworthiness of LLMs in RAG through grounded attributions and learning to refuse. In *International Conference on Learning Representations*.

Zhaochen Su, Jun Zhang, Xiaoye Qu, Tong Zhu, Yanshu Li, Jiashuo Sun, Juntao Li, Min Zhang, and Yu Cheng. 2024. ConflictBank: A benchmark for evaluating the influence of knowledge conflicts in LLMs. In *Advances in Neural Information Processing Systems (Datasets and Benchmarks Track)*, volume 37, pages 103242–103268.

Hexiang Tan, Fei Sun, Wanli Yang, Yuanzhuo Wang, Qi Cao, and Xueqi Cheng. 2024. Blinded by generated contexts: How language models merge generated and retrieved contexts when knowledge conflicts? In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 6207–6227, Bangkok, Thailand. Association for Computational Linguistics.

Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. 2023. Language models don't always say what they think: Unfaithful explanations in chain-of-thought prompting. In *Advances in Neural Information Processing Systems*.

Emiel van Miltenburg, Chris van der Lee, and Emiel Krahmer. 2021. Preregistering NLP research. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 613–623, Online. Association for Computational Linguistics.

Alexander Wan, Eric Wallace, and Dan Klein. 2024. What evidence do language models find convincing? In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 7468–7484.

Chenyu Wang, Yueyuan Li, Yingmin Liu, and Yang Shu. 2026. ConflictRAG: Detecting and resolving knowledge conflicts in retrieval augmented generation. *Preprint*, arXiv:2605.17301.

Fei Wang, Xingchen Wan, Ruoxi Sun, Jiefeng Chen, and Sercan O Arik. 2025a. Astute RAG: Overcoming imperfect retrieval augmentation and knowledge conflicts for large language models. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 30553–30571, Vienna, Austria. Association for Computational Linguistics.

Han Wang, Archiki Prasad, Elias Stengel-Eskin, and Mohit Bansal. 2025b. Retrieval-augmented generation with conflicting evidence. In *Second Conference on Language Modeling (COLM)*.

Yike Wang, Shangbin Feng, Heng Wang, Weijia Shi, Vidhisha Balachandran, Tianxing He, and Yulia Tsvetkov. 2024. Resolving knowledge conflicts in large language models. In *First Conference on Language Modeling*.

Zhepei Wei, Wei-Lin Chen, and Yu Meng. 2025. InstructRAG: Instructing retrieval-augmented generation via self-synthesized rationales. In *The Thirteenth International Conference on Learning Representations*.

Orion Weller, Aleem Khan, Nathaniel Weir, Dawn Lawrie, and Benjamin Van Durme. 2024. Defending against disinformation attacks in open-domain question answering. In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 402–417, St. Julian's, Malta. Association for Computational Linguistics.

Bingbing Wen, Jihan Yao, Shangbin Feng, Chenjun Xu, Yulia Tsvetkov, Bill Howe, and Lucy Lu Wang. 2025. Know your limits: A survey of abstention in large language models. *Transactions of the Association for Computational Linguistics*.

Kevin Wu, Eric Wu, and James Zou. 2024. ClashEval: Quantifying the tug-of-war between an LLM's internal prior and external evidence. In *Advances in Neural Information Processing Systems (Datasets and Benchmarks Track)*, volume 37, pages 33402–33422.

Chong Xiang, Tong Wu, Zexuan Zhong, David Wagner, Danqi Chen, and Prateek Mittal. 2024. Certifiably robust RAG against retrieval corruption. *Preprint*, arXiv:2405.15556.

Jian Xie, Kai Zhang, Jiangjie Chen, Renze Lou, and Yu Su. 2024. Adaptive chameleon or stubborn sloth: Revealing the behavior of large language models in knowledge conflicts. In *The Twelfth International Conference on Learning Representations (ICLR)*.

Rongwu Xu, Zehan Qi, Zhijiang Guo, Cunxiang Wang, Hongru Wang, Yue Zhang, and Wei Xu. 2024. Knowledge conflicts for LLMs: A survey. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 8541–8565, Miami, Florida, USA. Association for Computational Linguistics.

Shi-Qi Yan, Jia-Chen Gu, Yun Zhu, and Zhen-Hua Ling. 2024. Corrective retrieval augmented generation. *Preprint*, arXiv:2401.15884.

Hua Ye, Siyuan Chen, Ziqi Zhong, Canran Xiao, Haoliang Zhang, Yuhan Wu, and Fei Shen. 2026. Seeing through the conflict: Transparent knowledge conflict handling in retrieval-augmented generation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, pages 34423–34431.

Tian Yu, Shaolei Zhang, and Yang Feng. 2024. Truth-aware context selection: Mitigating hallucinations of large language models being misled by untruthful contexts. In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 10862–10884.

Yizhen Yuan, Rui Kong, Dongze Li, Yuanchun Li, and Yunxin Liu. 2026. Benchmarking LLM's capability in reasoning over conflicting web references. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 303–322, San Diego, California, United States. Association for Computational Linguistics.

Linda Zeng, Rithwik Gupta, Divij Motwani, Diji Yang, and Yi Zhang. 2025. Worse than zero-shot? A fact-checking dataset for evaluating the robustness of RAG against misleading retrievals. *arXiv preprint arXiv:2502.16101*.

Huichi Zhou, Kin-Hei Lee, Zhonghao Zhan, Yue Chen, Zhenhao Li, Zhaoyang Wang, Hamed Haddadi, and Emine Yilmaz. 2025. TrustRAG: Enhancing robustness and trustworthiness in retrieval-augmented generation. *arXiv preprint arXiv:2501.00879*.

Wei Zou, Runpeng Geng, Binghui Wang, and Jinyuan Jia. 2025. PoisonedRAG: Knowledge corruption attacks to retrieval-augmented generation of large language models. In *34th USENIX Security Symposium (USENIX Security 25)*, pages 3827–3844, Seattle, WA. USENIX Association.

A Prompts

All prompts below are copied verbatim from the code that produced the results.

Generation skeleton (default).

```
Answer the question using the passages below.
If the passages support more than one valid answer
(for example, different people or things share the
same name), give every valid answer.
If no passage answers the question, the answer is
"unknown".
{instruction}
{context}
{conflicts}
Question: {question}

Give a brief explanation, then end with one line in
exactly this format:
Final answer(s): <answer 1> | <answer 2> | ...
```

Minimal skeleton (no answer policy).

```
Answer the question using the passages below.
If no passage answers the question, the answer is
"unknown".
{instruction}
{context}
```

```
{conflicts}
Question: {question}

Give a brief explanation, then end with one line in
exactly this format:
Final answer(s): <your answer; if there is more than
one, separate them with |>
```

Single-answer skeleton.

```
Answer the question using the passages below. The
question has exactly one correct answer.
If no passage answers the question, the answer is
"unknown".
{instruction}
{context}
{conflicts}
Question: {question}

Give a brief explanation, then end with one line in
exactly this format:
Final answer(s): <your one answer>
```

Closed book (single-answer version).

```
Answer the question from your own knowledge. The
question has exactly one correct answer.
Question: {question}

Give a brief explanation, then end with one line in
exactly this format:
Final answer(s): <your one answer>
```

Instruction v1.

```
Some passages may contradict each other. Before
answering, check whether they do.
Ignore passages that appear to contain
misinformation, and if a disagreement reflects
genuinely different valid answers, report each one.
```

Instruction v2.

```
Some passages may contain false information. Before
answering, decide which passages are reliable by
checking them against each other.
Do not include an answer that is supported only by a
passage you judge to be unreliable.
```

Single call (inserted as the instruction).

```
Some passages may contain misinformation. First
decide which passages, if any, contain
misinformation, by checking them against each other
and against what you know, and write one line in
this format:
Suspect passages: <passage numbers, or none>
Then answer the question using only the other
passages.
```

Example: directional flag (passage 2 suspect).

```
Answer the question using the passages below.
If the passages support more than one valid answer
(for example, different people or things share the
same name), give every valid answer.
If no passage answers the question, the answer is
"unknown".

Passage 0 [POTENTIAL CONFLICT with passage(s): 2]:
<passage 0>

Passage 1:
<passage 1>

Passage 2 [LIKELY MISINFORMATION; conflicts with
passage(s): 0]:
<passage 2>

POTENTIAL CONFLICTS DETECTED:
- Passages 0 and 2: possible contradiction; passage
2 likely contains misinformation

Question: Who founded X?

Give a brief explanation, then end with one line in
exactly this format:
Final answer(s): <answer 1> | <answer 2> | ...
```

Example: one-wrong flag, single-answer skeleton.

```
Answer the question using the passages below. The
question has exactly one correct answer.
If no passage answers the question, the answer is
"unknown".

Passage 0 [CONFLICT with passage(s): 2; one side
contains misinformation]:
<passage 0>

Passage 1:
<passage 1>

Passage 2 [CONFLICT with passage(s): 0; one side
contains misinformation]:
<passage 2>

POTENTIAL CONFLICTS DETECTED:
- Passages 0 and 2: possible contradiction; exactly
one of the two contains misinformation (it is not
known which)

Question: Who founded X?

Give a brief explanation, then end with one line in
exactly this format:
Final answer(s): <your one answer>
```

Directional LLM detector.

```
You will see a question and {n} numbered passages
retrieved for it.
Some passages may contain misinformation: false
claims about the answer to the question.
Identify every pair of passages that CONTRADICT each
other about the answer to the question,
i.e. they cannot both be true about the same entity
at the same time. Passages that are about
different entities with the same name, give
complementary details, or describe different times
are NOT contradictions; note them as "ambiguity",
"complementary" or "temporal" instead.

For each contradiction, say which passage of the
pair is more likely the misinformation, judging
by agreement with the other passages and by
plausibility: its number, "both" if both are likely
false, or "unclear" if you cannot tell.

Question: {question}

{passages}

Return JSON only: {{"pairs": [{{"i": <int>, "j":
<int>, "relation":
"contradiction|ambiguity|complementary|temporal",
```

```
"misinformation": <int>|"both"|"unclear"}]}}
Use the passage numbers shown. "misinformation" is
needed only for contradictions. Return {{"pairs":
[]}} if there are none.
```

B Additional Results

Robustness to the metric and the labels. Table 21 repeats Table 3 with the bidirectional metric. Table 22 repeats the pre-registered tests with every suspect document treated as misinformation. Table 4 lists all pre-registered tests.

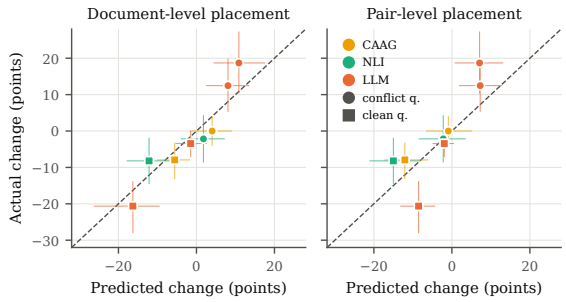


Figure 2: Effect of real detectors predicted from the synthetic grid (x, with the 95% bootstrap prediction interval) against the measured effect (y, with its 95% CI), for conflict (•) and clean (■) questions. Left: detectors located by their document-level behaviour; right: by pair precision, recall and direction accuracy.

The heuristic detector. CAAG scores a pair only if its MiniLM (Reimers and Gurevych, 2019) cosine similarity reaches $\theta_s = 0.65$. Its score is the sum of a negation and contrast asymmetry term, a named-entity mismatch term (differing dates, numbers, money or percentages) and a temporal-marker term, multiplied by the similarity. A pair is flagged if the score reaches $\theta_c = 0.35$. Its per-pair flag rate on the confirmatory split (6.22%) sets the rate of the random detector. CAAG is the detector we set out to study. As a neutral flag it did not beat random flags of the same volume, and as a removal rule it only deleted correct passages (−7.9 points on clean questions, 0.0 on conflict ones). Its false flags fall mostly on pairs that should both be reported: 134 of its 221 false flags on the hold-out are same-name ambiguity pairs, which its entity-mismatch signal flags because they mention different dates or numbers. We read this as a warning about lexical contradiction cues in general (Gururangan et al., 2018), rather than as a problem with a particular threshold. Retuning illustrates why F1 misleads. Over a 121-setting sweep, CAAG's best pair F1 is 29.2, above NLI's, at $\theta_s = 0.40$ and $\theta_c = 0.10$. But it gets there by flagging 43% of all pairs and 87% of clean questions, which our grid predicts would be harmful under removal (Appendix B).

Original CAAG format. Rendering CAAG's flags in its original format, with typed lines and confidence scores ("factual conflict (confidence: 0.41)"), changes accuracy by −0.9 points overall (−1.2 on conflict questions, −1.6 on clean ones). This is no different from the uniform format.

Pilot study (exploratory). On the 200-question pilot split, scored with the bidirectional metric, the same pattern appeared before any confirmatory run. Oracle neutral flags changed conflict-question accuracy by −10.5 points, directional flags by +11.8, and removing the suspect by +28.4. The pilots also showed that the original REMOVE strategy, which deletes an arbitrary side of each pair, removes correct documents. They are why we added directional formats, synthetic direction accuracy and document-level placement. All pilot runs are released.

Phase-0 audit of CAAG on NarrativeQA. CAAG was designed for long-document QA. On 100 NarrativeQA (Kočiský et al., 2018) questions from 50 stories (BM25 top-10 chunks from a single coherent story), it flags at least one pair in 54% of questions, 166 pairs in all. A single story rarely contradicts itself, so most of these flags are likely false; the flagged pairs are released for annotation. This is the setting where a false flag can only hurt.

C Further Results

C.1 Where the directional flag fails (exploratory)

After seeing the minimal-skeleton run, we split conflict questions by the number of gold answers. 105 of 139 have several (same-name ambiguity *and* misinformation), and a strict answer must list all of them. On the 34 single-answer conflict questions, the directional flag's estimate is positive in every setting, from +11.8 points (Qwen) to +46.1 (gpt-4o-mini), both with the default skeleton. On multi-answer questions it changes little except with gpt-4.1 (gpt-4o-mini +2.9, and −5.7 with the minimal skeleton; gpt-4.1 +19.0). One reading is that a warning about one passage makes the smaller models more conservative about the legitimate alternative answers too; we did not test

			Accuracy		Oracle pairs				LLM detector		Single	LLM dir.
Generator	Prompt	Seeds	Closed	Plain	Flag	One wrong	Dir. flag	Remove	Flag	Dir. flag	call	clean q.
gpt-4o-mini	single	3	10.0	51.1	**−18.7**	**−35.7**	**+30.1**	**+35.0**	**−18.2**	+8.4	+7.6	−4.7
gpt-4o-mini	multi	3	–	21.6	**−8.8**	–	**+27.9**	**+66.4**	**−8.4**	**+20.0**	–	−1.0
gpt-4o-mini	minimal	3	–	35.0	**−18.1**	–	**+38.4**	**+51.2**	–	–	–	–
gpt-4.1	single	1	20.0	64.7	**−8.0**	**−35.3**	**+26.7**	**+26.0**	–	−2.3	+7.3	−3.0
gpt-4.1	multi	1	–	13.0	**−6.3**	–	**+23.0**	**+78.3**	**−6.7**	**+17.0**	–	0.0
Qwen2.5-7B	single	1	4.3	50.3	−3.3	−5.0	**+25.3**	**+35.7**	−4.0	+7.3	–	−1.0
Qwen2.5-7B	multi	1	–	23.7	**−8.3**	–	**+17.7**	**+63.7**	−8.3	+6.7	–	−1.0
gpt-6-luna	single	3	16.9	53.6	**−9.4**	**−29.1**	**+36.4**	**+37.0**	**−9.3**	+8.6	−6.3	−1.3
gpt-6.1-terra	single	1	20.0	73.0	−2.3	−5.3	**+17.0**	**+18.3**	−2.7	−11.7	**−21.0**	0.0
Claude Haiku 5.5	single	3	11.8	66.8	+0.7	+6.6	**+23.3**	**+22.6**	+0.7	−4.1	+1.7	−5.0
Claude Sonnet 5.5	single	1	–	75.7	+3.0	+4.0	**+16.3**	–	–	−2.7	–	−1.0
gpt-4o-mini	forced	3	–	55.4	−0.2	**+7.6**	**+24.2**	–	–	+4.0	–	−2.7
gpt-6-luna	forced	3	–	64.6	+1.6	**+10.3**	**+26.4**	–	–	−2.4	–	−1.3
Claude Haiku 5.5	forced	3	–	69.3	**+3.9**	**+5.1**	**+21.1**	–	–	−4.7	–	−1.3
gpt-6.1-terra	forced	1	–	77.7	−0.3	**+6.7**	**+14.7**	–	–	**−14.3**	–	0.0
Majority misinformation: three wrong-side passages, one correct												
gpt-4o-mini	single	3	–	21.6	+1.1	−5.0	**+32.7**	**+65.1**	−2.2	**+5.7**	+6.8	−4.7
gpt-6-luna	single	3	–	29.8	**+6.6**	+0.3	**+60.4**	–	–	**+7.1**	–	−1.3
gpt-6.1-terra	single	1	–	26.0	−3.7	**+22.0**	**+66.3**	–	–	**+9.0**	–	0.0
Claude Haiku 5.5	single	3	–	31.0	+4.3	–	**+57.3**	–	–	**+6.8**	–	−5.0
RAGuard: real Reddit posts labelled supporting or misleading (294 conflict claims)												
gpt-4o-mini	single	3	64.9	36.7	+3.2	−5.7	**+9.9**	+5.9	–	+0.6	–	+1.7
gpt-6-luna	single	3	64.1	23.2	0.0	**−7.0**	**+8.8**	−1.0	–	+3.1	–	−1.3
Claude Haiku 5.5	single	3	65.1	24.5	+0.3	**+5.9**	**+12.7**	+2.0	–	**+5.0**	–	+1.3
gpt-6.1-terra	single	1	74.1	24.8	+2.4	+0.7	**+6.8**	−1.0	–	+2.0	–	+1.0
Claude Sonnet 5.5	single	1	–	27.2	+4.4	–	**+20.4**	–	–	**+7.1**	–	+10.0

Table 2: ConflictBank conflict questions (300; one correct answer each) and RAGuard conflict claims (294; answer true or false): accuracy (%, one-way strict) closed-book and with the passages but no flags, then the paired difference in points made by each format. *forced*: the one-answer prompt without the option to answer "unknown". *Prompt*: the answer policy: *single* says there is one answer, *multi* asks for every valid answer (our default skeleton), *minimal* states no policy. *One wrong*: the neutral flag plus "exactly one of the two contains misinformation". *Single call*: one call that names the suspect passages, then answers. Last column: the LLM directional flag's change on the 100 clean questions. The LLM detector is gpt-4o-mini for every generator. **Bold**: significant under the paper's single scheme.

this. Removal does not depend on the model acting on a warning, and it gains in both groups.

C.2 Real detectors

CAAG flags 34.1% of clean questions. Its neutral flags do not beat random flags at the same rate (H5; -0.4 vs. $+1.0$ points overall). The 90% CI of the difference, $[-3.1, +0.1]$, rules out a gain larger than $+0.1$ points but does not establish equivalence within ± 3 points (TOST). The directional-versus-neutral comparison with the LLM detector's pairs was registered as exploratory in the addendum. Under the bidirectional metric, the LLM detector's directional flag costs clean questions 6.1 points. With gpt-4.1 as the generator, its gain over all 300 questions ($+5.4$) is significant after Holm, but the 35 unlabelled questions contribute disproportionately to it (they gain 13.3 points); over the labelled questions it is $+4.4$ $[+0.6, +8.3]$. On ConflictBank the NLI cross-encoder flags no pair at all: ConflictBank passages are long, LLM-written and agree on everything but the one fact, and the truncated pair rarely contains it.

The single call with stronger models. Asked to name suspect passages before answering, the reasoning models and Claude Haiku distrust most of ConflictBank's generated passages and abstain, losing up to 51.0 points on clean questions (Table 27).

Reasoning models on RAMDocs. With no flags, gpt-6-luna and gpt-6.1-terra are strictly correct on few RAMDocs conflict questions because the default prompt asks for every valid answer and they list nearly every candidate: 2.8 and 3.0 answers on average, and 72% and 76% of their answers include the misinformation, while they recover 79% and 77% of the gold answers.

Detector parsing. The LLM detectors' JSON failed to parse on 1 of 2400 ConflictBank questions (a Claude Sonnet output), which then count

Condition	All	Lab.	Conflict	Clean	Tokens
No flags (plain)	29.9	31.3	22.1	41.5	911
Prompt only					
Instruction (v1)	+1.6	+0.8	−1.2	+2.9	948
Instruction (v2)	+3.1	+2.3	+5.3	−1.1	951
Oracle detector					
Oracle, flag	−3.1	−3.5	**−6.7**	0.0	980
Oracle, directional flag	+6.2	+7.0	**+13.4**	0.0	998
Oracle, remove suspect	**+12.8**	**+14.5**	**+27.6**	0.0	825
Real detectors (exploratory)					
Random, flag	+1.0	+1.0	+0.7	+1.3	949
CAAG, flag	−0.4	−0.8	−0.7	−0.8	942
CAAG, remove	−3.3	−3.8	0.0	−7.9	809
NLI, flag	+0.2	+0.3	−0.5	+1.1	1021
NLI, remove	−5.7	−5.0	−2.2	−8.2	592
LLM, flag	−1.8	−1.3	−4.1	+1.9	970
LLM, directional flag	+5.0	+4.9	+12.5	−3.4	983
LLM, remove suspect	+0.3	0.0	+18.7	**−20.6**	718
Reference					
MADAM-RAG	+5.7	+5.7	+3.6	+7.9	8196

Table 3: RAMDocs hold-out, gpt-4o-mini, 3 seeds, one-way strict accuracy (primary). *All*: 300 questions, including 35 whose conflict status the label rule leaves open; *Lab.*: the 265 labelled ones (139 conflict, 126 clean). First row: accuracy (%) with no flags. Other rows: paired difference from it in points; **bold** = significant under the paper's single scheme (Section 4: a registered test's Holm-adjusted p within its family, otherwise Holm over all exploratory comparisons). Real-detector rows were not pre-registered as tests. Tokens: mean prompt tokens per call (MADAM-RAG: per question, all calls).

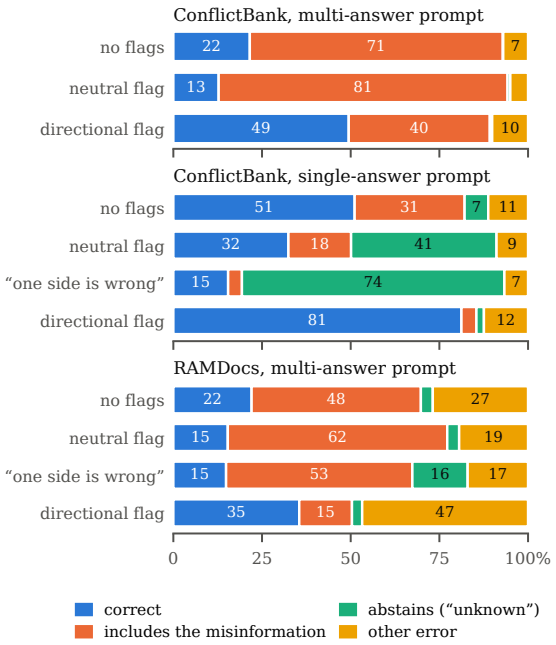


Figure 3: What a flag does to answers on conflict questions (gpt-4o-mini, oracle pairs; share of answers, %). The answer policy decides how a neutral flag fails: a model told to give every valid answer adds the misinformation, and one told there is one answer abstains. Only the directional flag raises the share of correct answers.

as flagging nothing.

C.3 A detector from another family

Replacing the gpt-4o-mini detector with Qwen2.5-7B (gpt-4o-mini still generates) gives a much less selective detector. It flags 96.8% of clean RAMDocs questions, but its direction accuracy is 89.0%. Its directional flag gains +10.6 points on conflict questions and 0.0 on clean ones, and its neutral flag −2.6. Removal again trades a conflict-question gain (+14.1) for a clean-question loss (−14.6). The conclusions do not depend on the detector's family.

C.4 CONFLICTS

On the 237 automatically scorable CONFLICTS questions (exploratory, 3 seeds; most have no conflict), no condition moves accuracy significantly (Table 24).

D Reproducibility

Runs. Every number in this paper is produced by one script from committed run outputs: per-call responses, token counts, detector outputs and configurations. The pre-registration was committed before any confirmatory run. Its deviations are listed below.

Deviations from the pre-registration.

1. *Primary metric.* The first draft of the pre-registration named the bidirectional metric as primary. Before the confirmatory run we switched to one-way matching, because it reproduces MADAM-RAG's published score on the pilot split (Section 4); bidirectional became the robustness metric. Every table reports both.

2. *Second generator.* The plan named an open-weights model. None was available before the confirmatory run, so gpt-4.1 was used (H9). We later added Qwen2.5-7B-Instruct, registered before its run. Because local inference takes about 3 s per call, the Qwen runs use one seed; we made that change before inspecting any Qwen result.

3. *CONFLICTS judge.* The plan called for a judge from a different model family, validated on about 100 human labels. Neither the judge nor the labels were available, so CONFLICTS is

	Pre-registered test	Estimate [95% CI]	p_{Holm}	Outcome	Robust
H1: detection quality					
H1	F1 llm_dir > nli	+2.6 [−2.4, +7.5]	0.311	not supported	–
H1	F1 nli > caag	+12.2 [+7.5, +16.9]	<0.001	supported	–
H1	direction accuracy llm_dir > 50% (cluster CI[r])	88.3% [83.3, 92.9]	<0.001	supported	–
H2: a correct neutral flag					
H2	oracle+flag < plain (conflict)	−6.7 [−10.8, −3.1]	<0.001	supported	yes
H3: what direction adds					
H3	oracle+flag_directional > plain (conflict)	+13.4 [+6.7, +20.4]	<0.001	supported	yes
H3	oracle+flag_directional > oracle+flag (conflict)	+20.1 [+13.9, +26.9]	<0.001	supported	yes
H3	oracle+remove_suspect > plain (conflict)	+27.6 [+20.6, +35.0]	<0.001	supported	yes
H4: false flags					
H4a	false-flag cells below plain on clean q., directional flag[r]	0 of 9 cells	Holm	not supported	yes
H4a	false-flag cells below plain on clean q., removal[r]	6 of 9 cells	Holm	supported	yes
H4b	harm grows with clean-q. exposure, Spearman over 18 cells[r]	$\rho = -0.39$ [−0.74, −0.12]	–	supported	yes
H4c	removal harms clean q. more than directional flag[r]	5 of 9 matched cells	Holm	count only	–
H5: the heuristic detector vs random					
H5	caag+flag vs random+flag (all)	−1.4 [−3.4, +0.3]	0.178	n.s.	yes
H5	equivalence within ±3 points (TOST)[r]	90% CI [−3.1, +0.1]	–	not equivalent	yes
H6: does the grid predict real detectors					
H6	actual effect inside the 95% interval of the document-level prediction[r]	7 of 8 (pair level: 6 of 8)	–	descriptive	–
H6	document-level error < pair-level error (paired, exact sign-flip)[r]	−1.2 [−3.6, +0.4]	0.414	not supported	yes
H7: prompt-only instructions					
H7	instruct vs plain (conflict)	−1.2 [−4.3, +1.7]	1.000	n.s.	yes
H7	instruct vs plain (clean)	+2.9 [0.0, +6.1]	0.280	n.s.	yes
H7	instruct_v2 vs plain (conflict)	+5.3 [+1.4, +9.6]	0.059	n.s.	yes
H7	instruct_v2 vs plain (clean)	−1.1 [−5.0, +2.6]	1.000	n.s.	yes
H9: gpt-4.1					
H9	gpt-4.1: oracle+flag < plain (conflict)	−8.6 [−13.2, −4.3]	<0.001	supported	yes
H9	gpt-4.1: oracle+flag_directional > plain (conflict)	+18.2 [+12.0, +24.7]	<0.001	supported	yes
H9	gpt-4.1: oracle+flag_directional > oracle+flag (conflict)	+26.9 [+20.1, +33.6]	<0.001	supported	yes
H9	gpt-4.1: oracle+remove_suspect > plain (conflict)	+30.2 [+22.5, +37.9]	<0.001	supported	yes
H8: genuine conflicts (WikiContradict)					
H8	WikiContradict: instruct vs plain	+8.7 [+5.4, +12.3]	<0.001	significant	–
H8	WikiContradict: instruct_v2 vs plain	−6.7 [−10.0, −3.7]	<0.001	significant	–
H8	WikiContradict: oracle+flag vs plain	+6.5 [+3.0, +10.0]	<0.001	supported	–
H8	WikiContradict: caag+flag vs plain	+0.7 [0.0, +1.6]	0.250	n.s.	–
H8	WikiContradict: nli+flag vs plain	+3.2 [+0.3, +6.1]	0.077	n.s.	–
H8	WikiContradict: llm_dir+flag vs plain	+6.2 [+2.8, +9.6]	<0.001	significant	–

Table 4: Pre-registered tests registered before the confirmatory run (RAMDocs hold-out unless stated; one-way strict accuracy; differences in points). Rows marked [r] are re-analyses registered later, before they were computed. Directional predictions are *supported* or *not supported*; tests with no predicted sign are *significant* or *n.s.* *Robust*: the same outcome under the bidirectional metric. p-values come from 10,000 sign-flip permutations, extended to 200,000 when none reaches the observed statistic, Holm-adjusted within the family; this is the significance used everywhere in the paper.

	Pre-registered test	Estimate [95% CI]	p_{Holm}	Outcome	Robust
H10: without the answer policy (minimal skeleton)					
H10	gpt-4o-mini: oracle+flag vs plain (conflict)	-5.3 [-9.4, -1.4]	0.042	significant	yes
H10	gpt-4o-mini: oracle+flag_directional > plain (conflict)	$+5.0$ [-1.7, $+11.8$]	0.314	not supported	yes
H10	gpt-4o-mini: oracle+flag_directional > oracle+flag (conflict)	$+10.3$ [$+4.8$, $+16.1$]	0.005	supported	yes
H10	gpt-4o-mini: oracle+remove_suspect > plain (conflict)	$+18.9$ [$+12.5$, $+25.9$]	<0.001	supported	yes
H10	gpt-4.1: oracle+flag vs plain (conflict)	-11.0 [-17.5, -4.6]	0.010	significant	yes
H10	gpt-4.1: oracle+flag_directional > plain (conflict)	$+12.7$ [$+5.8$, $+19.7$]	0.005	supported	yes
H10	gpt-4.1: oracle+flag_directional > oracle+flag (conflict)	$+23.7$ [$+17.0$, $+30.7$]	<0.001	supported	yes
H10	gpt-4.1: oracle+remove_suspect > plain (conflict)	$+14.9$ [$+8.6$, $+21.6$]	<0.001	supported	yes
H10	Qwen2.5-7B: oracle+flag vs plain (conflict)	-5.8 [-10.8, -1.4]	0.117	n.s.	yes
H10	Qwen2.5-7B: oracle+flag_directional > plain (conflict)	$+2.9$ [-3.6, $+9.4$]	0.509	not supported	yes
H10	Qwen2.5-7B: oracle+flag_directional > oracle+flag (conflict)	$+8.6$ [$+2.9$, $+14.4$]	0.041	supported	yes
H10	Qwen2.5-7B: oracle+remove_suspect > plain (conflict)	$+10.8$ [$+4.3$, $+17.3$]	0.017	supported	yes
H11: an open-weights generator					
H11	Qwen2.5-7B: oracle+flag < plain (conflict)	-5.8 [-10.8, -0.7]	0.114	not supported	yes
H11	Qwen2.5-7B: oracle+flag_directional > plain (conflict)	$+3.6$ [-0.7, $+8.6$]	0.224	not supported	yes
H11	Qwen2.5-7B: oracle+flag_directional > oracle+flag (conflict)	$+9.4$ [$+5.0$, $+14.4$]	0.001	supported	yes
H11	Qwen2.5-7B: oracle+remove_suspect > plain (conflict)	$+18.0$ [$+10.8$, $+25.2$]	<0.001	supported	yes

Table 5: Pre-registered tests added in the first addendum (§9; H10–H11), each committed before the runs it covers. Same conventions as Table 4. Settings name the benchmark and generator. ConflictBank rows from H15 on use the single-answer prompt unless marked *forced*; *CB majority* has three wrong-side passages per conflict question; RAGuard's misleading posts are real Reddit text.

	Pre-registered test	Estimate [95% CI]	p_{Holm}	Outcome	Robust
H12: ConflictBank					
H12	gpt-4o-mini: neutral flag vs plain (conflict)	−8.8 [−12.8, −4.9]	<0.001	significant	yes
H12	gpt-4o-mini: dir. flag > plain (conflict)	+27.9 [+22.8, +32.9]	<0.001	supported	yes
H12	gpt-4o-mini: dir. flag > neutral flag (conflict)	+36.7 [+31.6, +41.9]	<0.001	supported	yes
H12	gpt-4o-mini: removal > plain (conflict)	+66.4 [+61.1, +71.7]	<0.001	supported	yes
H12	gpt-4.1: neutral flag vs plain (conflict)	−6.3 [−9.7, −3.3]	<0.001	significant	yes
H12	gpt-4.1: dir. flag > plain (conflict)	+23.0 [+18.0, +28.3]	<0.001	supported	yes
H12	gpt-4.1: dir. flag > neutral flag (conflict)	+29.3 [+24.0, +34.7]	<0.001	supported	yes
H12	gpt-4.1: removal > plain (conflict)	+78.3 [+73.3, +83.0]	<0.001	supported	yes
H12	Qwen2.5-7B: neutral flag vs plain (conflict)	−8.3 [−12.3, −4.3]	<0.001	significant	yes
H12	Qwen2.5-7B: dir. flag > plain (conflict)	+17.7 [+12.0, +23.3]	<0.001	supported	yes
H12	Qwen2.5-7B: dir. flag > neutral flag (conflict)	+26.0 [+20.7, +31.3]	<0.001	supported	yes
H12	Qwen2.5-7B: removal > plain (conflict)	+63.7 [+58.3, +69.0]	<0.001	supported	yes
H13: neutral flag + resolution instruction					
H13	RAMDocs, gpt-4o-mini: neutral flag + instruction vs plain (conflict)	−2.4 [−6.2, +1.2]	0.743	n.s.	yes
H13	RAMDocs, gpt-4o-mini: dir. flag + instruction > neutral flag + instruction (conflict)	+15.6 [+9.6, +21.8]	<0.001	supported	yes
H13	RAMDocs, gpt-4.1: neutral flag + instruction vs plain (conflict)	+5.8 [−2.9, +14.4]	0.743	n.s.	yes
H13	RAMDocs, gpt-4.1: dir. flag + instruction > neutral flag + instruction (conflict)	+25.2 [+18.0, +32.4]	<0.001	supported	yes
H13	ConflictBank, gpt-4o-mini: neutral flag + instruction vs plain (conflict)	+1.6 [−2.8, +6.0]	0.743	n.s.	yes
H13	ConflictBank, gpt-4o-mini: dir. flag + instruction > neutral flag + instruction (conflict)	+62.2 [+57.1, +67.2]	<0.001	supported	yes
H14: published baselines					
H14	RAMDocs, gpt-4o-mini: LLM dir. flag vs astute multi (labelled)	+10.8 [+5.9, +16.0]	<0.001	significant	yes
H14	RAMDocs, gpt-4o-mini: LLM dir. flag vs trustrag multi (labelled)	+23.0 [+18.1, +28.3]	<0.001	significant	yes
H14	RAMDocs, gpt-4.1: LLM dir. flag vs astute multi (labelled)	+6.4 [+0.4, +12.5]	0.151	n.s.	no
H14	RAMDocs, gpt-4.1: LLM dir. flag vs trustrag multi (labelled)	+19.6 [+13.6, +25.7]	<0.001	significant	yes
H14	ConflictBank, gpt-4o-mini: LLM dir. flag vs astute (all questions)	−3.4 [−8.1, +1.3]	0.324	n.s.	yes
H14	ConflictBank, gpt-4o-mini: LLM dir. flag vs trustrag (all questions)	+0.8 [−3.9, +5.6]	0.761	n.s.	yes

Table 6: Pre-registered tests added in the second (§10; H12–H14), each committed before the runs it covers. Same conventions as Table 4. Settings name the benchmark and generator. ConflictBank rows from H15 on use the single-answer prompt unless marked *forced*; *CB majority* has three wrong-side passages per conflict question; RAGuard's misleading posts are real Reddit text.

	Pre-registered test	Estimate [95% CI]	p_{Holm}	Outcome	Robust
H15: ConflictBank, single-answer prompt					
H15	ConflictBank, gpt-4o-mini: neutral flag vs plain (conflict)	−18.7 [−23.4, −14.0]	<0.001	significant	yes
H15	ConflictBank, gpt-4o-mini: dir. flag > neutral flag (conflict)	+48.8 [+43.6, +54.1]	<0.001	supported	yes
H15	ConflictBank, gpt-4o-mini: dir. flag > plain (conflict)	+30.1 [+25.1, +35.2]	<0.001	supported	yes
H15	ConflictBank, gpt-4o-mini: LLM dir. flag > plain (all q.)	+5.2 [+1.0, +9.3]	0.065	not supported	no
H15	ConflictBank, gpt-4.1: neutral flag vs plain (conflict)	−8.0 [−12.7, −3.7]	0.006	significant	yes
H15	ConflictBank, gpt-4.1: dir. flag > neutral flag (conflict)	+34.7 [+29.0, +40.3]	<0.001	supported	yes
H15	ConflictBank, gpt-4.1: dir. flag > plain (conflict)	+26.7 [+21.7, +31.7]	<0.001	supported	yes
H15	ConflictBank, gpt-4.1: LLM dir. flag > plain (all q.)	−2.5 [−6.5, +1.5]	0.522	not supported	yes
H15	ConflictBank, Qwen2.5-7B: neutral flag vs plain (conflict)	−3.3 [−9.0, +2.0]	0.522	n.s.	yes
H15	ConflictBank, Qwen2.5-7B: dir. flag > neutral flag (conflict)	+28.7 [+23.3, +34.0]	<0.001	supported	yes
H15	ConflictBank, Qwen2.5-7B: dir. flag > plain (conflict)	+25.3 [+19.7, +31.0]	<0.001	supported	yes
H15	ConflictBank, Qwen2.5-7B: LLM dir. flag > plain (all q.)	+5.2 [+0.5, +10.0]	0.130	not supported	yes
H16: information vs direction (one-wrong flag)					
H16	ConflictBank, gpt-4o-mini: dir. flag > one-wrong flag (conflict)	+65.8 [+60.6, +70.9]	<0.001	supported	yes
H16	ConflictBank, gpt-4o-mini: one-wrong flag vs plain (conflict)	−35.7 [−41.1, −30.3]	<0.001	significant	yes
H16	ConflictBank, gpt-4.1: dir. flag > one-wrong flag (conflict)	+62.0 [+56.0, +67.7]	<0.001	supported	yes
H16	ConflictBank, gpt-4.1: one-wrong flag vs plain (conflict)	−35.3 [−41.3, −29.3]	<0.001	significant	yes
H16	RAMDocs, gpt-4o-mini: dir. flag > one-wrong flag (conflict)	+20.6 [+14.4, +27.3]	<0.001	supported	yes
H16	RAMDocs, gpt-4o-mini: one-wrong flag vs plain (conflict)	−7.2 [−12.2, −2.6]	0.006	significant	yes
H17: majority misinformation					
H17	CB majority, gpt-4o-mini: neutral flag vs plain (conflict)	+1.1 [−2.7, +4.9]	0.603	n.s.	yes
H17	CB majority, gpt-4o-mini: dir. flag > plain (conflict)	+32.7 [+27.3, +37.9]	<0.001	supported	yes
H17	CB majority, gpt-4o-mini: LLM dir. flag vs plain (conflict)	+5.7 [+1.4, +9.9]	0.017	significant	yes
H18: a single-call baseline					
H18	RAMDocs, gpt-4o-mini: LLM dir. flag vs self-identify (labelled)	−1.3 [−4.8, +2.1]	1.000	n.s.	yes
H18	ConflictBank, gpt-4o-mini: LLM dir. flag vs self-identify (all questions)	−0.2 [−4.2, +3.7]	1.000	n.s.	yes
H18	CB majority, gpt-4o-mini: LLM dir. flag vs self-identify (all questions)	−1.8 [−4.6, +1.1]	0.730	n.s.	yes
H19: current reasoning models					
H19	ConflictBank, gpt-6-luna: neutral flag vs plain (conflict)	−9.4 [−12.9, −6.0]	<0.001	significant	yes
H19	ConflictBank, gpt-6-luna: dir. flag > neutral flag (conflict)	+45.9 [+40.8, +50.9]	<0.001	supported	yes

Table 7: Pre-registered tests added in the third and fourth (§11–§12; H15–H19) (part 1 of 2), each committed before the runs it covers. Same conventions as Table 4. Settings name the benchmark and generator. ConflictBank rows from H15 on use the single-answer prompt unless marked *forced*; *CB majority* has three wrong-side passages per conflict question; RAGuard's misleading posts are real Reddit text.

	Pre-registered test	Estimate [95% CI]	p_{Holm}	Outcome	Robust
H19: current reasoning models					
H19	ConflictBank, gpt-6-luna: dir. flag > one-wrong flag (conflict)	+65.6 [+60.7, +70.1]	<0.001	supported	yes
H19	ConflictBank, gpt-6-luna: LLM dir. flag > plain (all q.)	+6.1 [+2.1, +10.0]	0.019	supported	yes
H19	ConflictBank, gpt-6.1-terra: neutral flag vs plain (conflict)	−2.3 [−6.7, +2.0]	0.717	n.s.	yes
H19	ConflictBank, gpt-6.1-terra: dir. flag > neutral flag (conflict)	+19.3 [+14.7, +24.0]	<0.001	supported	yes
H19	ConflictBank, gpt-6.1-terra: dir. flag > one-wrong flag (conflict)	+22.3 [+17.7, +27.3]	<0.001	supported	yes
H19	ConflictBank, gpt-6.1-terra: LLM dir. flag > plain (all q.)	−8.8 [−13.2, −4.2]	0.002	not supported	yes
H19	RAMDocs, gpt-6-luna: neutral flag vs plain (conflict)	−3.1 [−5.3, −1.4]	0.007	significant	yes
H19	RAMDocs, gpt-6-luna: dir. flag > neutral flag (conflict)	+42.2 [+34.5, +49.9]	<0.001	supported	yes
H19	RAMDocs, gpt-6.1-terra: neutral flag vs plain (conflict)	−0.7 [−2.9, +1.4]	1.000	n.s.	yes
H19	RAMDocs, gpt-6.1-terra: dir. flag > neutral flag (conflict)	+40.3 [+32.4, +48.2]	<0.001	supported	yes
H19	CB majority, gpt-6-luna: dir. flag > plain (conflict)	+60.4 [+55.3, +65.4]	<0.001	supported	yes
H19	CB majority, gpt-6-luna: LLM dir. flag vs plain (conflict)	+7.1 [+2.4, +11.8]	0.019	significant	yes
H19	CB majority, gpt-6.1-terra: dir. flag > plain (conflict)	+66.3 [+61.0, +71.7]	<0.001	supported	yes
H19	CB majority, gpt-6.1-terra: LLM dir. flag vs plain (conflict)	+9.0 [+3.3, +14.7]	0.019	significant	yes

Table 8: Pre-registered tests added in the third and fourth (§11–§12; H15–H19) (part 2 of 2), each committed before the runs it covers. Same conventions as Table 4. Settings name the benchmark and generator. ConflictBank rows from H15 on use the single-answer prompt unless marked *forced*; *CB majority* has three wrong-side passages per conflict question; RAGuard's misleading posts are real Reddit text.

	Pre-registered test	Estimate [95% CI]	p_{Holm}	Outcome	Robust
H20: forced choice (no "unknown")					
H20	ConflictBank forced, gpt-4o-mini: neutral flag vs plain (conflict)	-0.2 [-4.1, $+3.7$]	1.000	n.s.	yes
H20	ConflictBank forced, gpt-4o-mini: one-wrong flag vs plain (conflict)	$+7.6$ [$+2.7$, $+12.3$]	0.024	significant	yes
H20	ConflictBank forced, gpt-4o-mini: dir. flag > plain (conflict)	$+24.2$ [$+19.7$, $+29.0$]	<0.001	supported	yes
H20	ConflictBank forced, gpt-6-luna: neutral flag vs plain (conflict)	$+1.6$ [-1.6, $+4.8$]	1.000	n.s.	yes
H20	ConflictBank forced, gpt-6-luna: one-wrong flag vs plain (conflict)	$+10.3$ [$+5.6$, $+15.1$]	<0.001	significant	yes
H20	ConflictBank forced, gpt-6-luna: dir. flag > plain (conflict)	$+26.4$ [$+21.9$, $+31.0$]	<0.001	supported	yes
H20	ConflictBank forced, gpt-6.1-terra: neutral flag vs plain (conflict)	-0.3 [-4.0, $+3.3$]	1.000	n.s.	yes
H20	ConflictBank forced, gpt-6.1-terra: one-wrong flag vs plain (conflict)	$+6.7$ [$+2.3$, $+11.3$]	0.034	significant	yes
H20	ConflictBank forced, gpt-6.1-terra: dir. flag > plain (conflict)	$+14.7$ [$+10.3$, $+19.0$]	<0.001	supported	yes
H20	ConflictBank forced, claude-haiku-5-5: neutral flag vs plain (conflict)	$+3.9$ [$+1.2$, $+6.7$]	0.034	significant	yes
H20	ConflictBank forced, claude-haiku-5-5: one-wrong flag vs plain (conflict)	$+5.1$ [$+1.4$, $+8.8$]	0.034	significant	yes
H20	ConflictBank forced, claude-haiku-5-5: dir. flag > plain (conflict)	$+21.1$ [$+16.9$, $+25.4$]	<0.001	supported	yes
H22: pair topology					
H22	RAMDocs, gpt-4o-mini: one-wrong without pairs vs with pairs (conflict)	$+6.5$ [$+2.4$, $+10.8$]	0.015	significant	yes
H22	RAMDocs, gpt-6-luna: one-wrong without pairs vs with pairs (conflict)	$+10.1$ [$+5.8$, $+14.9$]	<0.001	significant	yes
H22	RAMDocs, gpt-6.1-terra: one-wrong without pairs vs with pairs (conflict)	-0.7 [-6.5, $+5.8$]	1.000	n.s.	yes
H22	CB majority, gpt-4o-mini: neutral without pairs vs with pairs (conflict)	-3.1 [-6.4, $+0.2$]	0.390	n.s.	yes
H22	CB majority, gpt-4o-mini: LLM dir. flag without pairs vs with pairs (conflict)	$+0.3$ [-2.1, $+2.8$]	1.000	n.s.	yes
H22	CB majority, gpt-6-luna: neutral without pairs vs with pairs (conflict)	-6.1 [-9.1, -3.1]	0.002	significant	yes
H22	CB majority, gpt-6-luna: LLM dir. flag without pairs vs with pairs (conflict)	$+0.8$ [-1.7, $+3.1$]	1.000	n.s.	yes
H22	CB majority, gpt-6.1-terra: neutral without pairs vs with pairs (conflict)	$+10.0$ [$+5.7$, $+14.3$]	0.002	significant	yes
H22	CB majority, gpt-6.1-terra: LLM dir. flag without pairs vs with pairs (conflict)	$+1.7$ [-1.3, $+4.7$]	1.000	n.s.	yes
H23: answer key					
H23	RAMDocs, gpt-4o-mini: dir. flag vs answer key (conflict)	-1.7 [-6.2, $+3.1$]	1.000	n.s.	yes
H23	RAMDocs, gpt-6-luna: dir. flag vs answer key (conflict)	-4.1 [-9.1, $+0.7$]	0.808	n.s.	yes
H23	RAMDocs, gpt-6.1-terra: dir. flag vs answer key (conflict)	-5.0 [-10.8, $+0.7$]	0.827	n.s.	yes
H23	ConflictBank, gpt-4o-mini: dir. flag vs answer key (conflict)	$+5.3$ [$+1.6$, $+9.1$]	0.059	n.s.	no
H23	ConflictBank, gpt-6-luna: dir. flag vs answer key (conflict)	$+0.1$ [-1.6, $+1.8$]	1.000	n.s.	yes

Table 9: Pre-registered tests added in the fifth and sixth (§13–§14; H20–H26) (part 1 of 5), each committed before the runs it covers. Same conventions as Table 4. Settings name the benchmark and generator. ConflictBank rows from H15 on use the single-answer prompt unless marked *forced*; *CB majority* has three wrong-side passages per conflict question; RAGuard's misleading posts are real Reddit text.

	Pre-registered test	Estimate [95% CI]	p_{Holm}	Outcome	Robust
H23: answer key					
H23	ConflictBank, gpt-6.1-terra: dir. flag vs answer key (conflict)	-1.7 [-3.3, -0.3]	0.515	n.s.	yes
H23	CB majority, gpt-4o-mini: dir. flag vs answer key (conflict)	$+3.2$ [-1.8, $+8.2$]	0.905	n.s.	yes
H23	CB majority, gpt-6-luna: dir. flag vs answer key (conflict)	$+0.4$ [-1.2, $+2.2$]	1.000	n.s.	yes
H23	CB majority, gpt-6.1-terra: dir. flag vs answer key (conflict)	$+2.3$ [$+0.3$, $+4.7$]	0.515	n.s.	yes
H25: Claude					
H25	ConflictBank, claude-haiku-5-5: neutral flag vs plain (conflict)	$+0.7$ [-2.2, $+3.4$]	1.000	n.s.	yes
H25	ConflictBank, claude-haiku-5-5: dir. flag > neutral flag (conflict)	$+22.7$ [$+18.6$, $+26.9$]	<0.001	supported	yes
H25	ConflictBank, claude-haiku-5-5: dir. flag > one-wrong flag (conflict)	$+16.8$ [$+13.0$, $+20.8$]	<0.001	supported	yes
H25	ConflictBank, claude-haiku-5-5: LLM dir. flag > plain (all q.)	-4.3 [-8.1, -0.7]	0.145	not supported	yes
H25	ConflictBank, claude-sonnet-5-5: neutral flag vs plain (conflict)	$+3.0$ [-0.3, $+6.3$]	0.634	n.s.	yes
H25	ConflictBank, claude-sonnet-5-5: dir. flag > neutral flag (conflict)	$+13.3$ [$+9.7$, $+17.3$]	<0.001	supported	yes
H25	ConflictBank, claude-sonnet-5-5: dir. flag > one-wrong flag (conflict)	$+12.3$ [$+8.7$, $+16.3$]	<0.001	supported	yes
H25	ConflictBank, claude-sonnet-5-5: LLM dir. flag > plain (all q.)	-2.2 [-5.8, $+1.0$]	0.993	not supported	yes
H25	RAMDocs, claude-haiku-5-5: neutral flag vs plain (conflict)	$+1.2$ [-2.6, $+5.3$]	1.000	n.s.	yes
H25	RAMDocs, claude-haiku-5-5: dir. flag > neutral flag (conflict)	$+27.1$ [$+20.1$, $+34.1$]	<0.001	supported	yes
H25	RAMDocs, claude-sonnet-5-5: neutral flag vs plain (conflict)	0.0 [-4.3, $+4.3$]	1.000	n.s.	yes
H25	RAMDocs, claude-sonnet-5-5: dir. flag > neutral flag (conflict)	$+14.4$ [$+7.9$, $+21.6$]	0.002	supported	yes
H25	CB majority, claude-haiku-5-5: dir. flag > plain (conflict)	$+57.3$ [$+52.0$, $+62.7$]	<0.001	supported	yes
H25	CB majority, claude-haiku-5-5: LLM dir. flag vs plain (conflict)	$+6.8$ [$+2.1$, $+11.4$]	0.043	significant	yes
H26: RAGuard (real misleading posts)					
H26	RAGuard, gpt-4o-mini: neutral flag vs plain (conflict)	$+3.2$ [$+0.1$, $+6.3$]	0.426	n.s.	yes
H26	RAGuard, gpt-4o-mini: dir. flag > plain (conflict)	$+9.9$ [$+5.6$, $+14.2$]	<0.001	supported	yes
H26	RAGuard, gpt-4o-mini: dir. flag > neutral flag (conflict)	$+6.7$ [$+2.9$, $+10.5$]	0.008	supported	yes
H26	RAGuard, gpt-4o-mini: one-wrong flag vs plain (conflict)	-5.7 [-10.2, -1.2]	0.161	n.s.	yes
H26	RAGuard, gpt-4o-mini: LLM dir. flag vs plain (conflict)	$+0.6$ [-2.8, $+4.0$]	1.000	n.s.	yes
H26	RAGuard, gpt-6-luna: neutral flag vs plain (conflict)	0.0 [-1.7, $+1.7$]	1.000	n.s.	yes
H26	RAGuard, gpt-6-luna: dir. flag > plain (conflict)	$+8.8$ [$+5.7$, $+12.1$]	<0.001	supported	yes
H26	RAGuard, gpt-6-luna: dir. flag > neutral flag (conflict)	$+8.8$ [$+5.7$, $+12.2$]	<0.001	supported	yes

Table 10: Pre-registered tests added in the fifth and sixth (§13–§14; H20–H26) (part 2 of 5), each committed before the runs it covers. Same conventions as Table 4. Settings name the benchmark and generator. ConflictBank rows from H15 on use the single-answer prompt unless marked *forced*; *CB majority* has three wrong-side passages per conflict question; RAGuard's misleading posts are real Reddit text.

	Pre-registered test	Estimate [95% CI]	p_{Holm}	Outcome	Robust
H26: RAGuard (real misleading posts)					
H26	RAGuard, gpt-6-luna: one-wrong flag vs plain (conflict)	−7.0 [−10.3, −3.9]	<0.001	significant	yes
H26	RAGuard, gpt-6-luna: LLM dir. flag vs plain (conflict)	+3.1 [−0.1, +6.2]	0.511	n.s.	yes
H26	RAGuard, gpt-6.1-terra: neutral flag vs plain (conflict)	+2.4 [0.0, +5.1]	0.731	n.s.	yes
H26	RAGuard, gpt-6.1-terra: dir. flag > plain (conflict)	+6.8 [+3.4, +10.5]	0.005	supported	yes
H26	RAGuard, gpt-6.1-terra: dir. flag > neutral flag (conflict)	+4.4 [+1.4, +7.8]	0.126	not supported	yes
H26	RAGuard, gpt-6.1-terra: one-wrong flag vs plain (conflict)	+0.7 [−3.4, +4.8]	1.000	n.s.	yes
H26	RAGuard, gpt-6.1-terra: LLM dir. flag vs plain (conflict)	+2.0 [−1.0, +5.1]	1.000	n.s.	yes
H26	RAGuard, claude-haiku-5-5: neutral flag vs plain (conflict)	+0.3 [−2.0, +2.7]	1.000	n.s.	yes
H26	RAGuard, claude-haiku-5-5: dir. flag > plain (conflict)	+12.7 [+9.2, +16.3]	<0.001	supported	yes
H26	RAGuard, claude-haiku-5-5: dir. flag > neutral flag (conflict)	+12.4 [+9.0, +15.9]	<0.001	supported	yes
H26	RAGuard, claude-haiku-5-5: one-wrong flag vs plain (conflict)	+5.9 [+2.9, +9.0]	0.005	significant	yes
H26	RAGuard, claude-haiku-5-5: LLM dir. flag vs plain (conflict)	+5.0 [+1.8, +8.2]	0.025	significant	yes
H26	RAGuard, claude-sonnet-5-5: neutral flag vs plain (conflict)	+4.4 [+0.7, +8.2]	0.321	n.s.	yes
H26	RAGuard, claude-sonnet-5-5: dir. flag > plain (conflict)	+20.4 [+15.6, +25.2]	<0.001	supported	yes
H26	RAGuard, claude-sonnet-5-5: dir. flag > neutral flag (conflict)	+16.0 [+11.6, +20.4]	<0.001	supported	yes
H26	RAGuard, claude-sonnet-5-5: LLM dir. flag vs plain (conflict)	+7.1 [+3.1, +11.2]	0.022	significant	yes
H21: risk and coverage					
H21	ConflictBank, gpt-4o-mini: neutral flag vs plain, misinformation rate (conflict)	−13.3 [−17.6, −9.1]	<0.001	significant	–
H21	ConflictBank, gpt-4o-mini: dir. flag vs plain, misinformation rate (conflict)	−26.8 [−31.6, −22.0]	<0.001	significant	–
H21	ConflictBank, gpt-6-luna: neutral flag vs plain, misinformation rate (conflict)	−6.9 [−9.9, −4.0]	<0.001	significant	–
H21	ConflictBank, gpt-6-luna: dir. flag vs plain, misinformation rate (conflict)	−24.3 [−28.6, −20.2]	<0.001	significant	–
H21	ConflictBank, gpt-6.1-terra: neutral flag vs plain, misinformation rate (conflict)	−2.3 [−6.0, +1.3]	1.000	n.s.	–
H21	ConflictBank, gpt-6.1-terra: dir. flag vs plain, misinformation rate (conflict)	−17.3 [−21.7, −13.0]	<0.001	significant	–
H21	ConflictBank, claude-haiku-5-5: neutral flag vs plain, misinformation rate (conflict)	+0.1 [−2.6, +2.8]	1.000	n.s.	–
H21	ConflictBank, claude-haiku-5-5: dir. flag vs plain, misinformation rate (conflict)	−21.8 [−26.1, −17.6]	<0.001	significant	–
H21	ConflictBank forced, gpt-4o-mini: neutral flag vs plain, misinformation rate (conflict)	+0.3 [−3.4, +4.2]	1.000	n.s.	–
H21	ConflictBank forced, gpt-4o-mini: dir. flag vs plain, misinformation rate (conflict)	−24.9 [−29.7, −20.1]	<0.001	significant	–

Table 11: Pre-registered tests added in the fifth and sixth (§13–§14; H20–H26) (part 3 of 5), each committed before the runs it covers. Same conventions as Table 4. Settings name the benchmark and generator. ConflictBank rows from H15 on use the single-answer prompt unless marked *forced*; *CB majority* has three wrong-side passages per conflict question; RAGuard's misleading posts are real Reddit text.

	Pre-registered test	Estimate [95% CI]	p_{Holm}	Outcome	Robust
H21: risk and coverage					
H21	ConflictBank forced, gpt-6-luna: neutral flag vs plain, misinformation rate (conflict)	−3.8 [−6.8, −0.9]	0.079	n.s.	–
H21	ConflictBank forced, gpt-6-luna: dir. flag vs plain, misinformation rate (conflict)	−25.7 [−30.0, −21.3]	<0.001	significant	–
H21	ConflictBank forced, gpt-6.1-terra: neutral flag vs plain, misinformation rate (conflict)	−1.3 [−5.0, +2.3]	1.000	n.s.	–
H21	ConflictBank forced, gpt-6.1-terra: dir. flag vs plain, misinformation rate (conflict)	−14.3 [−18.3, −10.3]	<0.001	significant	–
H21	ConflictBank forced, claude-haiku-5-5: neutral flag vs plain, misinformation rate (conflict)	−4.1 [−6.9, −1.3]	0.025	significant	–
H21	ConflictBank forced, claude-haiku-5-5: dir. flag vs plain, misinformation rate (conflict)	−20.9 [−25.1, −16.8]	<0.001	significant	–
H21	ConflictBank, gpt-4o-mini: neutral flag vs plain, abstention rate (conflict)	+34.3 [+29.3, +39.2]	<0.001	significant	–
H21	ConflictBank, gpt-4o-mini: dir. flag vs plain, abstention rate (conflict)	−4.6 [−7.2, −2.0]	0.013	significant	–
H21	ConflictBank, gpt-6-luna: neutral flag vs plain, abstention rate (conflict)	+17.0 [+13.6, +20.6]	<0.001	significant	–
H21	ConflictBank, gpt-6-luna: dir. flag vs plain, abstention rate (conflict)	−11.1 [−14.1, −8.2]	<0.001	significant	–
H21	ConflictBank, gpt-6.1-terra: neutral flag vs plain, abstention rate (conflict)	+3.0 [+0.7, +5.3]	0.281	n.s.	–
H21	ConflictBank, gpt-6.1-terra: dir. flag vs plain, abstention rate (conflict)	−1.3 [−2.7, −0.3]	1.000	n.s.	–
H21	ConflictBank, claude-haiku-5-5: neutral flag vs plain, abstention rate (conflict)	−0.3 [−0.9, +0.1]	1.000	n.s.	–
H21	ConflictBank, claude-haiku-5-5: dir. flag vs plain, abstention rate (conflict)	−0.1 [−0.9, +0.8]	1.000	n.s.	–
H21	ConflictBank forced, gpt-4o-mini: neutral flag vs plain, abstention rate (conflict)	+0.3 [−0.1, +0.9]	1.000	n.s.	–
H21	ConflictBank forced, gpt-4o-mini: dir. flag vs plain, abstention rate (conflict)	−0.1 [−0.3, 0.0]	1.000	n.s.	–
H21	ConflictBank forced, gpt-6-luna: neutral flag vs plain, abstention rate (conflict)	0.0 [0.0, 0.0]	1.000	n.s.	–
H21	ConflictBank forced, gpt-6-luna: dir. flag vs plain, abstention rate (conflict)	0.0 [0.0, 0.0]	1.000	n.s.	–
H21	ConflictBank forced, gpt-6.1-terra: neutral flag vs plain, abstention rate (conflict)	0.0 [0.0, 0.0]	1.000	n.s.	–
H21	ConflictBank forced, gpt-6.1-terra: dir. flag vs plain, abstention rate (conflict)	0.0 [0.0, 0.0]	1.000	n.s.	–
H21	ConflictBank forced, claude-haiku-5-5: neutral flag vs plain, abstention rate (conflict)	0.0 [0.0, 0.0]	1.000	n.s.	–
H21	ConflictBank forced, claude-haiku-5-5: dir. flag vs plain, abstention rate (conflict)	0.0 [0.0, 0.0]	1.000	n.s.	–
H21	ConflictBank, gpt-4o-mini: neutral flag vs plain, selective accuracy (conflict)	−1.0 [−4.9, +2.9]	1.000	n.s.	–
H21	ConflictBank, gpt-4o-mini: dir. flag vs plain, selective accuracy (conflict)	+28.2 [+23.0, +33.5]	<0.001	significant	–
H21	ConflictBank, gpt-6-luna: neutral flag vs plain, selective accuracy (conflict)	+0.3 [−3.3, +3.8]	1.000	n.s.	–
H21	ConflictBank, gpt-6-luna: dir. flag vs plain, selective accuracy (conflict)	+30.8 [+25.9, +35.6]	<0.001	significant	–

Table 12: Pre-registered tests added in the fifth and sixth (§13–§14; H20–H26) (part 4 of 5), each committed before the runs it covers. Same conventions as Table 4. Settings name the benchmark and generator. ConflictBank rows from H15 on use the single-answer prompt unless marked *forced*; *CB majority* has three wrong-side passages per conflict question; RAGuard's misleading posts are real Reddit text.

	Pre-registered test	Estimate [95% CI]	p_{Holm}	Outcome	Robust
H21: risk and coverage					
H21	ConflictBank, gpt-6.1-terra: neutral flag vs plain, selective accuracy (conflict)	0.0 [−3.9, +3.9]	1.000	n.s.	–
H21	ConflictBank, gpt-6.1-terra: dir. flag vs plain, selective accuracy (conflict)	+15.9 [+11.5, +20.3]	<0.001	significant	–
H21	ConflictBank, claude-haiku-5-5: neutral flag vs plain, selective accuracy (conflict)	+0.6 [−2.2, +3.3]	1.000	n.s.	–
H21	ConflictBank, claude-haiku-5-5: dir. flag vs plain, selective accuracy (conflict)	+23.2 [+18.9, +27.7]	<0.001	significant	–
H21	ConflictBank forced, gpt-4o-mini: neutral flag vs plain, selective accuracy (conflict)	−0.1 [−3.9, +3.8]	1.000	n.s.	–
H21	ConflictBank forced, gpt-4o-mini: dir. flag vs plain, selective accuracy (conflict)	+24.2 [+19.7, +29.0]	<0.001	significant	–
H21	ConflictBank forced, gpt-6-luna: neutral flag vs plain, selective accuracy (conflict)	+1.6 [−1.6, +4.8]	1.000	n.s.	–
H21	ConflictBank forced, gpt-6-luna: dir. flag vs plain, selective accuracy (conflict)	+26.4 [+21.9, +31.0]	<0.001	significant	–
H21	ConflictBank forced, gpt-6.1-terra: neutral flag vs plain, selective accuracy (conflict)	−0.3 [−4.0, +3.3]	1.000	n.s.	–
H21	ConflictBank forced, gpt-6.1-terra: dir. flag vs plain, selective accuracy (conflict)	+14.7 [+10.3, +19.0]	<0.001	significant	–
H21	ConflictBank forced, claude-haiku-5-5: neutral flag vs plain, selective accuracy (conflict)	+3.9 [+1.2, +6.7]	0.045	significant	–
H21	ConflictBank forced, claude-haiku-5-5: dir. flag vs plain, selective accuracy (conflict)	+21.1 [+16.9, +25.4]	<0.001	significant	–
H24: detector versus generator reliability (crossed)					
H24	gain rises with detector-minus-generator reliability, Spearman over 30 cells	$\rho = +0.64$	<0.001	supported	–
H24	per-question slope of gain on reliability gap, within cells	+47.6 [+43.2, +51.9]	<0.001	supported	–

Table 13: Pre-registered tests added in the fifth and sixth (§13–§14; H20–H26) (part 5 of 5), each committed before the runs it covers. Same conventions as Table 4. Settings name the benchmark and generator. ConflictBank rows from H15 on use the single-answer prompt unless marked *forced*; *CB majority* has three wrong-side passages per conflict question; RAGuard's misleading posts are real Reddit text.

	Pre-registered test	Estimate [95% CI]	p_{Holm}	Outcome	Robust
H27: the detector model's answer as a hint					
H27	Qwen2.5-7B, detector gpt-4.1: LLM dir. flag vs hint (conflict)	+8.7 [+2.7, +14.3]	0.093	n.s.	no
H27	Qwen2.5-7B, detector gpt-4o-mini: LLM dir. flag vs hint (conflict)	+4.3 [−1.7, +10.3]	1.000	n.s.	yes
H27	Qwen2.5-7B, detector gpt-6-luna: LLM dir. flag vs hint (conflict)	+4.3 [−1.3, +10.0]	1.000	n.s.	yes
H27	Qwen2.5-7B, detector gpt-6.1-terra: LLM dir. flag vs hint (conflict)	−4.0 [−9.7, +1.7]	1.000	n.s.	yes
H27	claude-haiku-5-5, detector claude-haiku-5-5: LLM dir. flag vs hint (conflict)	+11.1 [+7.2, +15.1]	<0.001	significant	yes
H27	claude-haiku-5-5, detector claude-sonnet-5-5: LLM dir. flag vs hint (conflict)	+14.6 [+9.6, +19.7]	<0.001	significant	yes
H27	claude-haiku-5-5, detector gpt-4.1: LLM dir. flag vs hint (conflict)	+14.1 [+9.8, +18.6]	<0.001	significant	yes
H27	claude-haiku-5-5, detector gpt-4o-mini: LLM dir. flag vs hint (conflict)	−1.4 [−6.0, +3.2]	1.000	n.s.	yes
H27	claude-haiku-5-5, detector gpt-6-luna: LLM dir. flag vs hint (conflict)	+10.3 [+6.6, +14.1]	<0.001	significant	yes
H27	claude-haiku-5-5, detector gpt-6.1-terra: LLM dir. flag vs hint (conflict)	+11.0 [+7.1, +15.0]	<0.001	significant	yes
H27	gpt-4.1, detector gpt-4.1: LLM dir. flag vs hint (conflict)	+13.0 [+8.0, +18.0]	<0.001	significant	yes
H27	gpt-4.1, detector gpt-4o-mini: LLM dir. flag vs hint (conflict)	+0.7 [−4.7, +6.0]	1.000	n.s.	yes
H27	gpt-4.1, detector gpt-6-luna: LLM dir. flag vs hint (conflict)	+9.7 [+4.7, +14.7]	0.006	significant	yes
H27	gpt-4.1, detector gpt-6.1-terra: LLM dir. flag vs hint (conflict)	+1.7 [−3.0, +6.3]	1.000	n.s.	yes
H27	gpt-4o-mini, detector claude-haiku-5-5: LLM dir. flag vs hint (conflict)	+6.2 [+1.1, +11.4]	0.279	n.s.	yes
H27	gpt-4o-mini, detector claude-sonnet-5-5: LLM dir. flag vs hint (conflict)	+5.8 [+0.2, +11.3]	0.485	n.s.	yes
H27	gpt-4o-mini, detector gpt-4.1: LLM dir. flag vs hint (conflict)	+11.8 [+6.6, +17.0]	0.005	significant	yes
H27	gpt-4o-mini, detector gpt-4o-mini: LLM dir. flag vs hint (conflict)	+7.1 [+1.8, +12.6]	0.180	n.s.	yes
H27	gpt-4o-mini, detector gpt-6-luna: LLM dir. flag vs hint (conflict)	+8.1 [+2.9, +13.2]	0.037	significant	yes
H27	gpt-4o-mini, detector gpt-6.1-terra: LLM dir. flag vs hint (conflict)	+0.1 [−4.9, +5.2]	1.000	n.s.	yes
H27	gpt-6-luna, detector claude-haiku-5-5: LLM dir. flag vs hint (conflict)	+4.9 [−0.4, +10.3]	0.785	n.s.	yes
H27	gpt-6-luna, detector claude-sonnet-5-5: LLM dir. flag vs hint (conflict)	+3.9 [−0.8, +8.6]	0.992	n.s.	yes
H27	gpt-6-luna, detector gpt-4.1: LLM dir. flag vs hint (conflict)	+13.7 [+8.4, +18.9]	<0.001	significant	yes
H27	gpt-6-luna, detector gpt-4o-mini: LLM dir. flag vs hint (conflict)	+6.0 [+0.6, +11.7]	0.485	n.s.	yes
H27	gpt-6-luna, detector gpt-6-luna: LLM dir. flag vs hint (conflict)	+8.2 [+3.7, +12.8]	0.005	significant	yes
H27	gpt-6-luna, detector gpt-6.1-terra: LLM dir. flag vs hint (conflict)	−5.8 [−10.6, −1.0]	0.279	n.s.	yes

Table 14: Pre-registered tests added in the seventh (§15; H27–H29) (part 1 of 5), each committed before the runs it covers. Same conventions as Table 4. Settings name the benchmark and generator. ConflictBank rows from H15 on use the single-answer prompt unless marked *forced*; *CB majority* has three wrong-side passages per conflict question; RAGuard's misleading posts are real Reddit text.

	Pre-registered test	Estimate [95% CI]	p_{Holm}	Outcome	Robust
H27: the detector model's answer as a hint					
H27	gpt-6.1-terra, detector gpt-4.1: LLM dir. flag vs hint (conflict)	+10.7 [+5.0, +16.0]	0.008	significant	yes
H27	gpt-6.1-terra, detector gpt-4o-mini: LLM dir. flag vs hint (conflict)	−6.3 [−12.0, −0.7]	0.485	n.s.	yes
H27	gpt-6.1-terra, detector gpt-6-luna: LLM dir. flag vs hint (conflict)	+9.7 [+4.7, +14.7]	0.005	significant	yes
H27	gpt-6.1-terra, detector gpt-6.1-terra: LLM dir. flag vs hint (conflict)	+9.7 [+5.0, +14.3]	0.008	significant	yes
H27	Qwen2.5-7B, detector gpt-4.1: hint vs plain (conflict)	+7.3 [+2.3, +12.3]	0.071	n.s.	yes
H27	Qwen2.5-7B, detector gpt-4o-mini: hint vs plain (conflict)	+3.0 [−2.0, +8.0]	1.000	n.s.	yes
H27	Qwen2.5-7B, detector gpt-6-luna: hint vs plain (conflict)	+7.7 [+3.0, +12.7]	0.038	significant	no
H27	Qwen2.5-7B, detector gpt-6.1-terra: hint vs plain (conflict)	+12.0 [+7.0, +17.0]	<0.001	significant	yes
H27	claude-haiku-5-5, detector claude-haiku-5-5: hint vs plain (conflict)	−5.4 [−8.8, −2.1]	0.032	significant	yes
H27	claude-haiku-5-5, detector claude-sonnet-5-5: hint vs plain (conflict)	−6.3 [−9.4, −3.2]	0.004	significant	yes
H27	claude-haiku-5-5, detector gpt-4.1: hint vs plain (conflict)	−3.1 [−5.8, −0.4]	0.341	n.s.	yes
H27	claude-haiku-5-5, detector gpt-4o-mini: hint vs plain (conflict)	−2.7 [−5.6, +0.1]	0.681	n.s.	yes
H27	claude-haiku-5-5, detector gpt-6-luna: hint vs plain (conflict)	−5.3 [−8.2, −2.6]	0.011	significant	yes
H27	claude-haiku-5-5, detector gpt-6.1-terra: hint vs plain (conflict)	−6.9 [−9.9, −4.0]	<0.001	significant	yes
H27	gpt-4.1, detector gpt-4.1: hint vs plain (conflict)	−1.3 [−3.3, +0.3]	1.000	n.s.	yes
H27	gpt-4.1, detector gpt-4o-mini: hint vs plain (conflict)	−3.0 [−6.7, +0.3]	1.000	n.s.	yes
H27	gpt-4.1, detector gpt-6-luna: hint vs plain (conflict)	−3.3 [−7.3, +0.3]	1.000	n.s.	yes
H27	gpt-4.1, detector gpt-6.1-terra: hint vs plain (conflict)	+3.3 [−0.7, +7.0]	1.000	n.s.	yes
H27	gpt-4o-mini, detector claude-haiku-5-5: hint vs plain (conflict)	+6.8 [+2.8, +10.9]	0.027	significant	yes
H27	gpt-4o-mini, detector claude-sonnet-5-5: hint vs plain (conflict)	+10.1 [+6.4, +13.8]	<0.001	significant	yes
H27	gpt-4o-mini, detector gpt-4.1: hint vs plain (conflict)	+6.8 [+3.7, +10.0]	<0.001	significant	yes
H27	gpt-4o-mini, detector gpt-4o-mini: hint vs plain (conflict)	+1.3 [−0.4, +3.2]	1.000	n.s.	yes
H27	gpt-4o-mini, detector gpt-6-luna: hint vs plain (conflict)	+5.6 [+2.2, +9.0]	0.036	significant	no
H27	gpt-4o-mini, detector gpt-6.1-terra: hint vs plain (conflict)	+9.4 [+5.9, +13.1]	<0.001	significant	yes
H27	gpt-6-luna, detector claude-haiku-5-5: hint vs plain (conflict)	+9.4 [+5.7, +13.2]	<0.001	significant	yes
H27	gpt-6-luna, detector claude-sonnet-5-5: hint vs plain (conflict)	+14.2 [+10.4, +18.2]	<0.001	significant	yes

Table 15: Pre-registered tests added in the seventh (§15; H27–H29) (part 2 of 5), each committed before the runs it covers. Same conventions as Table 4. Settings name the benchmark and generator. ConflictBank rows from H15 on use the single-answer prompt unless marked *forced*; *CB majority* has three wrong-side passages per conflict question; RAGuard's misleading posts are real Reddit text.

	Pre-registered test	Estimate [95% CI]	p_{Holm}	Outcome	Robust
H27: the detector model's answer as a hint					
H27	gpt-6-luna, detector gpt-4.1: hint vs plain (conflict)	+7.7 [+4.2, +11.2]	<0.001	significant	yes
H27	gpt-6-luna, detector gpt-4o-mini: hint vs plain (conflict)	+2.6 [−1.3, +6.4]	1.000	n.s.	yes
H27	gpt-6-luna, detector gpt-6-luna: hint vs plain (conflict)	+2.9 [+0.2, +5.6]	0.462	n.s.	yes
H27	gpt-6-luna, detector gpt-6.1-terra: hint vs plain (conflict)	+13.1 [+9.4, +16.9]	<0.001	significant	yes
H27	gpt-6.1-terra, detector gpt-4.1: hint vs plain (conflict)	−3.3 [−7.7, +1.0]	1.000	n.s.	yes
H27	gpt-6.1-terra, detector gpt-4o-mini: hint vs plain (conflict)	−5.3 [−9.3, −1.3]	0.194	n.s.	yes
H27	gpt-6.1-terra, detector gpt-6-luna: hint vs plain (conflict)	−7.0 [−11.0, −3.3]	0.006	significant	yes
H27	gpt-6.1-terra, detector gpt-6.1-terra: hint vs plain (conflict)	−7.0 [−11.0, −3.0]	0.027	significant	yes
H27	Qwen2.5-7B, detector gpt-4.1: cell vs detector model alone (conflict)	+1.7 [−4.0, +7.3]	1.000	n.s.	yes
H27	Qwen2.5-7B, detector gpt-4o-mini: cell vs detector model alone (conflict)	+6.6 [+0.8, +12.3]	0.312	n.s.	yes
H27	Qwen2.5-7B, detector gpt-6-luna: cell vs detector model alone (conflict)	+8.8 [+3.6, +14.0]	0.021	significant	yes
H27	Qwen2.5-7B, detector gpt-6.1-terra: cell vs detector model alone (conflict)	−14.7 [−20.3, −9.0]	<0.001	significant	yes
H27	claude-haiku-5-5, detector claude-haiku-5-5: cell vs detector model alone (conflict)	+5.7 [+2.1, +9.3]	0.027	significant	no
H27	claude-haiku-5-5, detector claude-sonnet-5-5: cell vs detector model alone (conflict)	−0.7 [−4.9, +3.6]	1.000	n.s.	yes
H27	claude-haiku-5-5, detector gpt-4.1: cell vs detector model alone (conflict)	+13.1 [+8.1, +18.0]	<0.001	significant	yes
H27	claude-haiku-5-5, detector gpt-4o-mini: cell vs detector model alone (conflict)	+11.6 [+5.4, +17.4]	0.009	significant	yes
H27	claude-haiku-5-5, detector gpt-6-luna: cell vs detector model alone (conflict)	+18.2 [+13.3, +23.1]	<0.001	significant	yes
H27	claude-haiku-5-5, detector gpt-6.1-terra: cell vs detector model alone (conflict)	−2.1 [−7.6, +3.3]	1.000	n.s.	yes
H27	gpt-4.1, detector gpt-4.1: cell vs detector model alone (conflict)	+11.7 [+6.7, +16.7]	<0.001	significant	yes
H27	gpt-4.1, detector gpt-4o-mini: cell vs detector model alone (conflict)	+11.2 [+5.2, +17.1]	0.006	significant	yes
H27	gpt-4.1, detector gpt-6-luna: cell vs detector model alone (conflict)	+17.4 [+12.4, +22.3]	<0.001	significant	yes
H27	gpt-4.1, detector gpt-6.1-terra: cell vs detector model alone (conflict)	−3.3 [−8.7, +2.0]	1.000	n.s.	yes
H27	gpt-4o-mini, detector claude-haiku-5-5: cell vs detector model alone (conflict)	−2.7 [−7.7, +2.3]	1.000	n.s.	yes
H27	gpt-4o-mini, detector claude-sonnet-5-5: cell vs detector model alone (conflict)	−8.7 [−13.7, −3.7]	0.018	significant	no
H27	gpt-4o-mini, detector gpt-4.1: cell vs detector model alone (conflict)	+5.0 [−0.2, +10.1]	0.671	n.s.	yes
H27	gpt-4o-mini, detector gpt-4o-mini: cell vs detector model alone (conflict)	+8.4 [+3.2, +13.8]	0.021	significant	yes

Table 16: Pre-registered tests added in the seventh (§15; H27–H29) (part 3 of 5), each committed before the runs it covers. Same conventions as Table 4. Settings name the benchmark and generator. ConflictBank rows from H15 on use the single-answer prompt unless marked *forced*; *CB majority* has three wrong-side passages per conflict question; RAGuard's misleading posts are real Reddit text.

	Pre-registered test	Estimate [95% CI]	p_{Holm}	Outcome	Robust
H27: the detector model's answer as a hint					
H27	gpt-4o-mini, detector gpt-6-luna: cell vs detector model alone (conflict)	+11.2 [+5.8, +16.7]	0.004	significant	yes
H27	gpt-4o-mini, detector gpt-6.1-terra: cell vs detector model alone (conflict)	−12.3 [−18.0, −6.7]	<0.001	significant	yes
H27	gpt-6-luna, detector claude-haiku-5-5: cell vs detector model alone (conflict)	+1.1 [−3.8, +5.8]	1.000	n.s.	yes
H27	gpt-6-luna, detector claude-sonnet-5-5: cell vs detector model alone (conflict)	−4.0 [−8.3, +0.3]	0.671	n.s.	yes
H27	gpt-6-luna, detector gpt-4.1: cell vs detector model alone (conflict)	+10.2 [+4.8, +15.6]	0.009	significant	yes
H27	gpt-6-luna, detector gpt-4o-mini: cell vs detector model alone (conflict)	+11.0 [+4.9, +17.0]	0.011	significant	yes
H27	gpt-6-luna, detector gpt-6-luna: cell vs detector model alone (conflict)	+11.1 [+7.3, +15.0]	<0.001	significant	yes
H27	gpt-6-luna, detector gpt-6.1-terra: cell vs detector model alone (conflict)	−12.1 [−17.2, −6.9]	<0.001	significant	yes
H27	gpt-6.1-terra, detector gpt-4.1: cell vs detector model alone (conflict)	+15.7 [+10.7, +20.7]	<0.001	significant	yes
H27	gpt-6.1-terra, detector gpt-4o-mini: cell vs detector model alone (conflict)	+10.2 [+4.2, +16.2]	0.018	significant	yes
H27	gpt-6.1-terra, detector gpt-6-luna: cell vs detector model alone (conflict)	+22.1 [+17.2, +27.0]	<0.001	significant	yes
H27	gpt-6.1-terra, detector gpt-6.1-terra: cell vs detector model alone (conflict)	+2.7 [−0.3, +5.7]	0.925	n.s.	yes
H28: LLM flags on genuine conflicts					
H28	WikiContradict, gpt-4o-mini: LLM dir. flag vs plain	−6.5 [−10.3, −2.8]	0.001	significant	–
H28	WikiContradict, gpt-4o-mini: LLM removal vs plain	−9.1 [−13.3, −5.1]	<0.001	significant	–
H28	WikiContradict, gpt-6-luna: LLM dir. flag vs plain	−29.0 [−34.5, −23.6]	<0.001	significant	–
H28	WikiContradict, gpt-6-luna: LLM removal vs plain	−30.0 [−35.6, −24.5]	<0.001	significant	–
H28	WikiContradict, claude-haiku-5-5: LLM dir. flag vs plain	−27.7 [−33.2, −22.3]	<0.001	significant	–
H28	WikiContradict, claude-haiku-5-5: LLM removal vs plain	−28.2 [−33.6, −22.9]	<0.001	significant	–
H28	RAMDocs ambiguity, gpt-4o-mini: LLM dir. flag vs plain	−4.0 [−8.4, 0.0]	0.100	n.s.	–
H28	RAMDocs ambiguity, gpt-4o-mini: LLM removal vs plain	−21.2 [−29.6, −13.1]	<0.001	significant	–
H28	RAMDocs ambiguity, gpt-4o-mini: LLM dir. flag vs plain, gold recall	−3.6 [−6.8, −0.7]	0.024	significant	–
H28	RAMDocs ambiguity, gpt-4o-mini: LLM removal vs plain, gold recall	−17.4 [−22.4, −12.4]	<0.001	significant	–
H28	RAMDocs ambiguity, gpt-6-luna: LLM dir. flag vs plain	−14.8 [−21.5, −8.8]	<0.001	significant	–
H28	RAMDocs ambiguity, gpt-6-luna: LLM removal vs plain	−29.3 [−38.0, −20.9]	<0.001	significant	–
H28	RAMDocs ambiguity, gpt-6-luna: LLM dir. flag vs plain, gold recall	−8.9 [−12.4, −5.7]	<0.001	significant	–
H28	RAMDocs ambiguity, gpt-6-luna: LLM removal vs plain, gold recall	−20.3 [−25.0, −15.7]	<0.001	significant	–

Table 17: Pre-registered tests added in the seventh (§15; H27–H29) (part 4 of 5), each committed before the runs it covers. Same conventions as Table 4. Settings name the benchmark and generator. ConflictBank rows from H15 on use the single-answer prompt unless marked *forced*; *CB majority* has three wrong-side passages per conflict question; RAGuard's misleading posts are real Reddit text.

	Pre-registered test	Estimate [95% CI]	p_{Holm}	Outcome	Robust
H28: LLM flags on genuine conflicts					
H28	RAMDocs ambiguity, claude-haiku-5-5: LLM dir. flag vs plain	-5.1 [-10.1, -0.3]	0.068	n.s.	–
H28	RAMDocs ambiguity, claude-haiku-5-5: LLM removal vs plain	-25.6 [-34.0, -17.8]	<0.001	significant	–
H28	RAMDocs ambiguity, claude-haiku-5-5: LLM dir. flag vs plain, gold recall	-2.7 [-5.8, $+0.2$]	0.089	n.s.	–
H28	RAMDocs ambiguity, claude-haiku-5-5: LLM removal vs plain, gold recall	-16.6 [-21.2, -12.0]	<0.001	significant	–
H28	RAMDocs ambiguity, gpt-6.1-terra: LLM dir. flag vs plain	-1.0 [-6.1, $+3.0$]	1.000	n.s.	–
H28	RAMDocs ambiguity, gpt-6.1-terra: LLM removal vs plain	-27.3 [-36.4, -18.2]	<0.001	significant	–
H28	RAMDocs ambiguity, gpt-6.1-terra: LLM dir. flag vs plain, gold recall	0.0 [-2.4, $+2.2$]	1.000	n.s.	–
H28	RAMDocs ambiguity, gpt-6.1-terra: LLM removal vs plain, gold recall	-18.5 [-23.6, -13.5]	<0.001	significant	–
H29: artefact probe of detector direction					
H29	gpt-4o-mini detector, reversed order: direction accuracy vs original	-1.3 [-7.3, $+4.7$]	1.000	n.s.	–
H29	gpt-4o-mini detector, paraphrased: direction accuracy vs original	$+1.0$ [-2.7, $+4.7$]	1.000	n.s.	–
H29	gpt-6-luna detector, reversed order: direction accuracy vs original	-1.7 [-6.0, $+2.7$]	1.000	n.s.	–
H29	gpt-6-luna detector, paraphrased: direction accuracy vs original	-1.3 [-5.7, $+3.0$]	1.000	n.s.	–
H29	gpt-4o-mini detector, reversed order: gpt-4o-mini's gain vs original	-1.9 [-8.4, $+4.7$]	1.000	n.s.	–
H29	gpt-4o-mini detector, paraphrased: gpt-4o-mini's gain vs original	$+1.7$ [-3.1, $+6.3$]	1.000	n.s.	–
H29	gpt-6-luna detector, reversed order: gpt-4o-mini's gain vs original	-4.6 [-8.9, -0.2]	0.196	n.s.	–
H29	gpt-6-luna detector, paraphrased: gpt-4o-mini's gain vs original	-1.2 [-5.1, $+2.7$]	1.000	n.s.	–

Table 18: Pre-registered tests added in the seventh (§15; H27–H29) (part 5 of 5), each committed before the runs it covers. Same conventions as Table 4. Settings name the benchmark and generator. ConflictBank rows from H15 on use the single-answer prompt unless marked *forced*; *CB majority* has three wrong-side passages per conflict question; RAGuard's misleading posts are real Reddit text.

		Plain	Oracle, conflict q.			LLM, directional flag		LLM, remove suspect	
Generator	Skeleton	conflict	flag	dir. flag	remove	conflict	clean	conflict	clean
gpt-4o-mini	default	22.1	**−6.7**	**+13.4**	**+27.6**	+12.5	−3.4	+18.7	**−20.6**
gpt-4o-mini	minimal	20.4	**−5.3**	+5.0	**+18.9**	+3.6	−10.3	+11.0	**−19.6**
gpt-4.1	default	21.8	**−8.6**	**+18.2**	**+30.2**	+8.4	0.0	**+20.6**	**−24.3**
gpt-4.1	minimal	33.6	**−11.0**	**+12.7**	**+14.9**	–	–	–	–
Qwen2.5-7B	default	19.4	−5.8	+3.6	**+18.0**	+2.9	−7.1	+7.9	**−18.3**
Qwen2.5-7B	minimal	17.3	−5.8	+2.9	**+10.8**	–	–	–	–
gpt-6-luna‡	default	14.4	**−3.1**	**+39.1**	**+38.8**	**+29.7**	**−11.6**	–	–
gpt-6.1-terra‡	default	10.8	−0.7	**+39.6**	**+39.6**	**+38.8**	−0.8	–	–

Table 19: The main effects across generators and prompt skeletons (RAMDocs hold-out, 3 seeds, one-way strict accuracy). *Plain*: accuracy (%) on conflict questions with no flags. Other columns: paired difference from plain in points. *Default* skeleton tells the model to give every valid answer; *minimal* drops that policy (Appendix A). **Bold**: significant under the paper's single scheme (Section 4: a registered test's Holm-adjusted p within its family, otherwise Holm over all exploratory comparisons). The LLM detector is gpt-4o-mini for every generator. Qwen2.5-7B-Instruct runs locally (4-bit), seed 0. ‡: current reasoning models at their default temperature (gpt-6-luna 3 seeds, gpt-6.1-terra seed 0).

	Pair level			Clean-q.	Direc-	Removal, per question		Latency	Cost
Detector	P	R	F1	flagged	tion	misinfo gone	correct lost	(s)	($/1k q.)
CAAG (heuristic)	21.7	7.9	11.6	34.1	–	9.4	31.7	0.23	0.000
NLI cross-encoder	20.1	29.3	23.8	61.9	–	44.6	52.4	1.18	0.000
LLM (gpt-4o-mini)	37.2	20.5	26.4	76.2	88.3	61.2	65.9	0.80	0.181

Table 20: Detectors on the RAMDocs hold-out split (300 questions; percentages). *Clean-q. flagged*: share of conflict-free questions with at least one flag. *Direction*: share of correctly flagged contradictions whose suspect side is the misinformation. *Misinfo gone*: conflict questions whose misinformation documents are all removed when removing the suspect (or, without direction, the second passage). *Correct lost*: clean questions that lose at least one correct document. Latency is per question; CAAG and NLI run locally.

scored only on items with a reference answer, with one seed, and reported as exploratory.

4. *H4.* The pre-registration counted 54 false-flag cells. On clean questions, cells that differ only in direction accuracy render identical prompts (their scores differ by exactly 0), so they are one comparison: 18 in all, 9 per format. We report H4a per format. The cell-permutation p-value for the dose-response correlation assumed exchangeable cells; the cells are nested, so we replace it with a question-level bootstrap CI. No threshold was set for H4c, so we report a count.

5. *H5.* A non-significant difference is not equivalence. We add a TOST with a ± 3-point margin, chosen in the addendum after the confirmatory run, so it is post hoc.

6. *H6.* The first draft tested whether the grid prediction lies inside the actual effect's CI, which is lenient. We now report the pre-registered test (actual effect inside a bootstrap prediction interval) and a paired test of the two spaces' errors.

7. *Placement modes.* For H6, the detectors without direction were run under the label `remove_suspect` as well as `remove`. Both render identical prompts, so these were served entirely from the response cache.

8. *Significance.* Earlier drafts used different corrections in different tables. We now apply one scheme throughout (Section 4), which makes some exploratory differences that were bold in earlier drafts non-significant.

9. *Single- and multi-answer split.* The split of conflict questions by number of gold answers (Appendix C.1) was chosen after seeing the minimal-skeleton run and is exploratory.

10. *Second registration.* We later registered H12–H14 (ConflictBank, the neutral flag with a resolution instruction, and the published baselines), the Qwen2.5-7B detector and the ConflictBank sampling rule. The gpt-4.1 runs in it use seed 0 only, to stay within budget. One ConflictBank gpt-4o-mini run stopped on a transient API error and was resumed from the response cache; no row was re-sampled.

Condition	All	Lab.	Conflict	Clean	Tokens
No flags (plain)	35.2	37.2	25.9	49.7	911
Prompt only					
Instruction (v1)	+1.3	+0.5	−1.0	+2.1	948
Instruction (v2)	+2.2	+1.4	+3.6	−1.1	951
Oracle detector					
Oracle, flag	−3.6	−4.0	**−7.7**	0.0	980
Oracle, directional flag	+7.0	**+7.9**	**+15.1**	0.0	998
Oracle, remove suspect	**+15.4**	**+17.5**	**+33.3**	0.0	825
Real detectors (exploratory)					
Random, flag	+0.4	+0.4	−0.2	+1.1	949
CAAG, flag	−0.9	−1.4	−1.9	−0.8	942
CAAG, remove	−4.0	−4.5	−1.0	−8.5	809
NLI, flag	+0.2	0.0	−1.4	+1.6	1021
NLI, remove	−8.0	−7.9	−3.1	−13.2	592
LLM, flag	−3.4	−3.4	−7.2	+0.8	970
LLM, directional flag	+4.7	+4.5	+14.1	−6.1	983
LLM, remove suspect	−2.1	−2.6	+18.9	**−26.5**	718
Reference					
MADAM-RAG	+5.7	+5.2	+3.6	+6.9	8196

Table 21: RAMDocs hold-out, gpt-4o-mini, 3 seeds, bidirectional strict accuracy (robustness). *All*: 300 questions, including 35 whose conflict status the label rule leaves open; *Lab.*: the 265 labelled ones (139 conflict, 126 clean). First row: accuracy (%) with no flags. Other rows: paired difference from it in points; **bold** = significant under the paper's single scheme (Section 4: a registered test's Holm-adjusted p within its family, otherwise Holm over all exploratory comparisons). Real-detector rows were not pre-registered as tests. Tokens: mean prompt tokens per call (MADAM-RAG: per question, all calls).

11. *Third and fourth registrations.* H15–H18 (§11) and H19 (§12) were committed before their runs. To stay within budget, §12 dropped two §11 runs before they started: gpt-4.1 under majority misinformation (replaced by gpt-6.1-terra in H17) and gpt-4.1 seeds 1–2, so H14 with gpt-4.1 stays on seed 0 as first registered and every gpt-4.1 baseline comparison is seed-0 matched.

12. *Reasoning models.* gpt-6-luna and gpt-6.1-terra accept only their default temperature and require `max_completion_tokens`; we used their defaults with a 4,096-token budget. 25 responses from the two models (almost all closed-book) were empty because reasoning used the whole budget; they are scored as wrong.

13. *Fifth and sixth registrations.* H20–H24 (§13) and H25–H26 with Claude and RAGuard (§14) were committed before their runs. The planned `cross_mj` runs completed within budget. Gemini was planned but not run, because no API key was available. Claude rejects a temperature and JSON mode on the endpoint we used; it ran at its default temperature, and as a detector its JSON is parsed by the detector's regex fallback. The answer parser unwraps paired tags (`<answer>X</answer>`), which no earlier result contained.

14. *Permutation resolution.* With 10,000 permutations the smallest p-value is 10^{-4}, which cannot survive Holm over 636 exploratory comparisons. After the §14 runs, and before writing them up, we extended the permutations to 200,000 whenever none of the first 10,000 reached the observed statistic. This changes only p-values at their floor, and only lowers them.

15. *H24 additions.* The cross-cell break-even point, the count of positive cells and the mechanism score (credit for naming the false side where the generator errs, charge for marking a correct passage where it is right) were defined after seeing the H24 cells and are exploratory.

16. *Seventh registration.* H27–H29 and the re-analyses of §15 were committed before their runs. To finish within the session, the §15 runs were started in parallel tracks rather than strictly in priority order; spend never approached the $195 limit, so no run was dropped. The question-level model adds one dependency (statsmodels 0.14.5). H21(c) is now the ratio of correct to answered answers; the registered per-question mean averaged the two conditions over different questions (those not abstained in every seed), which is why an earlier draft's selective accuracies did not match Table 1. Both are in the released statistics.

17. *Run interruptions.* Several runs were stopped and resumed (a transient API error; reordering the run queue). Resumption keeps completed rows and re-requests only missing ones; no row was re-sampled.

The addenda (H10–H29 and the re-analyses above) and the Qwen seed change were each committed to the repository before the runs they cover. The repository's commit history is the only timestamp; we did not use a third-party registry (van Miltenburg et al., 2021).

RAMDocs labels. The released document types are imperfect: 163 documents typed *correct* contain a wrong answer, some as injected misinformation, others incidentally. We leave pairs that touch them out of pair-level metrics, and questions

	Test (labels: *contains*)	One-way [95% CI]	p_{Holm}	Bidir.	p_{Holm}
H2	oracle+flag < plain (conflict)	−4.6 [−8.1, −1.5]	0.006	−5.2	0.003
H3	oracle+flag_directional > plain (conflict)	+11.6 [+5.6, +17.3]	<0.001	+13.3	<0.001
H3	oracle+flag_directional > oracle+flag (conflict)	+16.2 [+10.8, +21.8]	<0.001	+18.5	<0.001
H3	oracle+remove_suspect > plain (conflict)	+25.8 [+19.5, +32.2]	<0.001	+30.3	<0.001
H4	harmful cells, clean q. (dir. flag; removal); ρ	0/9; 8/9; $\rho = -0.08$ [-0.56, -0.02]			

Table 22: Label robustness: pre-registered tests with every suspect document treated as misinformation (173 conflict and 127 clean hold-out questions).

CAAG setting	θ_s	θ_c	P	R	F1	Clean FPR
Default	0.65	0.35	21.7	7.9	11.6	34.1
Best pair F1	0.40	0.10	20.7	49.6	29.2	87.3
Best question F1	0.55	0.15	22.2	28.8	25.1	60.3

Signals (defaults)	Pair F1	Flag rate	Clean FPR
all signals	11.6	6.2	34.1
negation only	0.6	0.4	4.0
entity only	8.6	4.3	23.0
temporal only	0.0	0.0	0.0

Table 23: CAAG on the RAMDocs hold-out (%). Top: thresholds from a sweep over 121 settings ($\theta_s \in [0.40, 0.90]$, $\theta_c \in [0.10, 0.60]$). Bottom: each signal alone.

Condition	All	Conflict	No conflict
No flags (plain)	65.8	46.9	73.6
Instruction (v2)	+0.6	+3.9	−0.8
CAAG, flag	+1.1	+2.4	+0.6
NLI, flag	−0.3	−0.5	−0.2
LLM, flag	+2.7	+5.3	+1.6
LLM, directional flag	+0.3	+1.4	−0.2

Table 24: CONFLICTS, 237 automatically scorable questions (69 conflict), gpt-4o-mini, 3 seeds; exploratory. Contains-match accuracy (%) and differences in points; bold would mark Holm-adjusted $p < 0.05$ within this table, and no difference reaches it. CAAG: flags 94.2% of conflict and 92.9% of clean questions; NLI: flags 52.2% of conflict and 43.5% of clean questions; LLM: flags 98.6% of conflict and 83.9% of clean questions.

Condition	All	Lab.	Conflict	Clean
No flags (plain)	32.7	34.5	21.8	48.4
Instruction (v2)	+5.8	+5.2	+12.2	−2.6
Oracle, flag	−4.0	−4.5	**−8.6**	0.0
Oracle, directional flag	**+8.4**	**+9.6**	**+18.2**	0.0
Oracle, remove suspect	**+14.0**	**+15.8**	**+30.2**	0.0
LLM, directional flag	+5.4	+4.4	+8.4	0.0
LLM, remove suspect	+1.0	−0.8	**+20.6**	**−24.3**
NLI, remove	−7.0	−7.3	+0.2	**−15.6**

Table 25: Replication with gpt-4.1 as generator (RAMDocs hold-out, 3 seeds, one-way strict accuracy; detector fixed to gpt-4o-mini). Same conventions as Table 3.

whose only sign of conflict is such a document (55 of 500; 35 of the confirmatory split) count only in "All". Treating them as misinformation is a robustness check.

Power. With 139 RAMDocs conflict questions, the minimum detectable effect at 80% power is 4–12 points (median 5.8; Card et al., 2020).

Metric choice. With the one-way metric, our port of MADAM-RAG scores 28.5% on the pilot split and 32.8% on all 500 questions (published: 28.00, on an undisclosed subset); with the bidirectional metric it scores 38.4% on all 500. Because the published subset and scorer are unknown, we do not claim to reproduce the published number. Normalisation lower-cases and removes punctuation and articles.

WikiContradict and CONFLICTS. A WikiContradict answer is correct if both of the dataset's answer strings (one per passage) match a final answer under the bidirectional rule. From CONFLICTS we use the 237 questions that have a reference answer (outdated-information, misinformation and no-conflict items); an answer is correct if it contains the reference answer.

Baselines. MADAM-RAG is ported unchanged to our client: three debate rounds, temperature 0, 128 tokens per call. Astute RAG has no official code. We use the prompts in its appendix verbatim with its default settings (t=1 consolidation round, m=1 generated passage, temperature 0, 1,024 output tokens). A generated passage that says "I don't know" is dropped, and the answer is the text in its `<ANSWER>` tag. TrustRAG is ported from its repository (commit `11dcea0`, MIT licence): SimCSE CLS embeddings (`sup-simcse-bert-base-uncased`), a ROUGE-L gate of 0.25, 2-means clustering and a cosine threshold of 0.88, followed by its three prompts verbatim at temperature 0.1. Its system

Condition	gpt-4o-mini	gpt-4.1
No flags (plain)	14.0	29.0
Instruction (v1)	+**8.7**	–
Instruction (v2)	−**6.7**	–
Oracle, flag	+**6.5**	+**19.0**
CAAG, flag	+0.7	–
NLI, flag	+3.2	–
LLM, flag	+**6.2**	–

Table 26: WikiContradict (253 genuine conflicts between two Wikipedia passages; the reference answer is both sides), 3 seeds (gpt-4.1: the oracle flag only). First row: share of answers giving both sides (%); other rows: difference in points, **bold** = significant under the paper's single scheme (H8 is a registered family; the gpt-4.1 column is exploratory).

message is not sent, because our client sends user messages only. The *multi* variants change one sentence of the final step to ask for every valid answer (Astute RAG) or append our final-answer line (TrustRAG); the unmodified variants are also reported.

ConflictBank. We use the QA file of the dataset's Hugging Face release (revision `92866a0`). Questions are sampled round-robin over relations with seed 0; distractors are other subjects' correct passages. The sampling script, the item ids and the seed are released; the data are not redistributed, because the release states a licence in its README but ships no licence file.

RAGuard. We use the Hugging Face release `UCSC-IRKM/RAGuard` (revision `dae344a2`, MIT licence on the card). A post is usable if it has at least 15 words after removing the "[Link Post]" placeholder and is not a duplicate within its claim; posts are cut to 300 words and shuffled by claim id. Supporting × misleading pairs are the contradictions and misleading posts the misinformation side. The build script and item ids are released.

Cost and compute. All API calls in this paper, pilots included, cost $185.14. Prices were checked on 2026-10-27 (per million input, cached-input and output tokens: gpt-4o-mini $0.15, $0.075 and $0.60; gpt-4.1 $2.00, $0.50 and $8.00; gpt-6-luna $0.10, $0.01 and $0.50; gpt-6.1-terra $1.00, $0.05 and $5.00; Claude Haiku 5.5 $0.10, $0.01 and $0.50; Claude Sonnet 5.5 $2.00, $0.10 and $10.00). CAAG and NLI ran on a laptop CPU (0.2 s and 1.2 s per question).

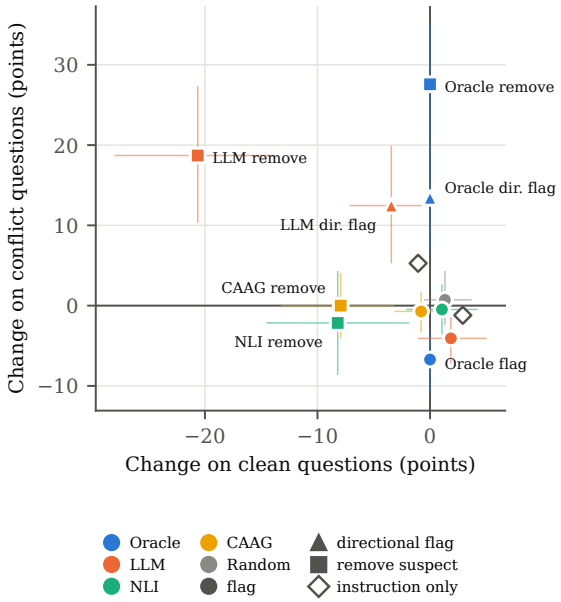


Figure 4: Change in one-way strict accuracy on RAMDocs conflict questions (y) against clean questions (x), relative to no flags (gpt-4o-mini, default skeleton, confirmatory split, 95% CIs). Colour: detector; shape: format. The oracle sits on the y-axis because it never flags a clean question. Without direction, removal deletes the second passage of each pair.

E Synthetic Detectors and the Grid

Construction. Each synthetic detector (Section 3.3) is built from the gold labels. Each gold contradiction pair is flagged with probability r. False flags are drawn from non-contradicting pairs at a single per-pair rate, set over the whole split so that expected precision is p. That rate is dataset-wide, so conflict-free questions receive false flags too. A true flag names the gold misinformation side with probability d, and otherwise the other side. A false flag has no misinformation side, so its suspect is picked uniformly. Each pair gets one uniform draw for flagging and one for direction, and every cell reuses them. Cells are therefore nested: raising r or d, or lowering p, only adds flags or corrections. This keeps comparisons across cells paired.

Results. Figure 5 replaces the oracle with synthetic detectors of known quality. Gains rise with direction accuracy in both formats. Only the cells with perfect direction are significant on all questions (directional flag +6, removal +13 points at perfect precision), but neighbouring cells at direction accuracy 0.75 have similar point estimates on conflict questions, so this is a dose-response, not

Method	RAMDocs, gpt-4o-mini			RAMDocs, gpt-4.1		ConflictBank		CB majority		
	Lab.	Conflict	Clean	Lab.	Conflict	Conflict	Clean	Conflict	Clean	Calls
No flags (plain), accuracy	31.3	22.1	41.5	32.8†	19.4†	51.1	89.0	21.6	89.0	1
Instruction (v2) only	+2.3	+5.3	−1.1	–	+12.2	–	–	–	–	1
Oracle, flag + instr.	−1.8	−2.4	−1.1	+1.5†	+5.8†	–	–	–	–	1
Oracle, dir. flag + instr.	+6.4	+13.2	−1.1	**+14.7**†	**+30.9**†	–	–	–	–	1
MADAM-RAG	+5.7	+3.6	+7.9	–	–	–	–	–	–	9×tok.
Astute RAG	**−15.6**	−8.4	**−23.5**	–	–	−2.1	−9.7	+1.4	−9.7	2
Astute RAG (multi)	−5.9	−1.2	−11.1	+2.3†	+13.7†	–	–	–	–	2
TrustRAG	**−17.5**	−11.3	**−24.3**	–	–	−9.7	−4.0	+6.8	−4.0	3
TrustRAG (multi)	**−18.1**	−12.2	**−24.6**	−10.9†	0.0†	–	–	–	–	3
Single call (name suspects, then answer)	+6.2	+12.0	−0.3	–	–	+7.6	−1.0	+6.8	−1.0	1
LLM detector, flag + instr.	−2.0	−2.2	−1.9	–	–	–	–	–	–	1+1
LLM detector, dir. flag + instr.	−3.5	+8.4	**−16.7**	–	–	–	–	–	–	1+1
LLM detector, dir. flag	+4.9	+12.5	−3.4	+8.7†	+14.4†	+8.4	−4.7	**+5.7**	−4.7	1+1
LLM detector, remove suspect	0.0	+18.7	**−20.6**	–	**+20.6**	+7.9	−7.0	+7.0	−7.0	1+1

Table 27: Resolution instructions, published baselines and a single-call baseline. First row: accuracy (%, one-way strict) with no flags; other rows: paired difference from it in points. RAMDocs uses the multi-answer prompt (it has same-name questions with several answers); ConflictBank and its majority variant (three wrong-side passages, one correct) use the single-answer prompt and gpt-4o-mini. *Lab.*: labelled RAMDocs questions. 3 seeds, except †: gpt-4.1 seed 0 only, relative to the seed-0 no-flag accuracy in the first row (other gpt-4.1 cells are relative to its 3-seed no-flag accuracy, Table 25). *instr.*: the v2 reliability instruction. *(multi)*: the baseline's final step asks for every valid answer. *Single call*: one call that names suspect passages, then answers. Calls per question; *1+1*: one detector call and one generator call. **Bold**: significant under the paper's single scheme.

a threshold: with 36 cells under Holm, the grid has little power to separate adjacent levels. Precision matters mainly for removal. On clean questions a cell's prompt does not depend on direction accuracy, so the false-flag cells reduce to 9 distinct comparisons per format. Removal is significantly harmful in 6 of them, the directional flag in 0 (H4a): false flags make the model delete correct passages, which costs clean questions 15 points at precision 0.5 and 26 at 0.25, while synthetic false directional flags cost at most two points. The harm tracks the share of clean questions that lose a correct document (Spearman ρ =-0.39 over 18 distinct cells, 95% CI [−0.74, −0.12] from a question-level bootstrap; H4b). Removal hurt clean questions significantly more than the directional flag in 5 of 9 matched cells (H4c; we set no threshold, so we report the count). Removal is the more powerful and the more fragile interface: its best cell gains twice the directional flag's on all questions, but below precision 0.75 it is the only format that does net harm.

Placing real detectors on the grid. Can the grid predict a real detector's effect from a few summary numbers? We compare two ways of summarising a detector. *Pair-level*: its pair precision, recall and direction accuracy. *Document-level*: what it does to each subset's documents in a given format. For conflict questions, that is the share whose misinformation documents are all marked and the share that have a correct document marked; for clean questions, the share that have a correct document marked. Here "marked" means removed (Remove) or headed as misinformation (Directional flag). The prediction for a subset is the inverse-distance-weighted mean effect of the four nearest grid cells in that space.

Does the grid predict real detectors? If the grid only described synthetic detectors, it would be of limited use. Figure 2 (Appendix B) tests whether it predicts real ones (H6). As preregistered, we bootstrap the whole prediction (cell effects, every detector's location and the interpolation) over questions, and ask whether the actual effect falls inside the 95% interval. Locating detectors by what they do to documents (the share of conflict questions whose misinformation is marked, and the share of questions that have a correct document marked) puts 7 of 8 effects inside; locating them by pair precision, recall and direction accuracy puts 6 inside. Document-level errors are smaller on average (4.1 vs. 5.3 points), but the paired difference, −1.2 [−3.6, +0.4] (exact sign-flip $p = 0.41$), is not significant over 8 cases. Both spaces miss the same case, the LLM detector's removal gain on conflict questions: one

		Detector					Accuracy with			Flag minus	
Detector	Generator	Right	Wrong	Gen. acc.	Score	Gain	flag	hint	det. alone	hint	det. alone
gpt-4o-mini	gpt-6.1-terra	67.0	30.3	73.0	−4.7	**−11.7**	61.3	67.7	51.1	−6.3	**+10.2**
gpt-4o-mini	Claude Haiku	67.0	30.3	66.8	+1.9	−4.1	62.7	64.1	51.1	−1.4	**+11.6**
gpt-6.1-terra	gpt-6.1-terra	29.7	4.3	73.0	+2.7	+2.7	75.7	66.0	73.0	**+9.7**	+2.7
gpt-4o-mini	gpt-4.1	67.0	30.3	64.7	+4.0	−2.3	62.3	61.7	51.1	+0.7	**+11.2**
gpt-6.1-terra	Claude Haiku	29.7	4.3	66.8	+4.3	**+4.1**	70.9	59.9	73.0	**+11.0**	−2.1
gpt-6.1-terra	gpt-4.1	29.7	4.3	64.7	+4.7	**+5.0**	69.7	68.0	73.0	+1.7	−3.3
gpt-6-luna	gpt-6.1-terra	45.7	6.7	73.0	+5.3	+2.7	75.7	66.0	53.6	**+9.7**	**+22.1**
gpt-6-luna	Claude Haiku	45.7	6.7	66.8	+7.1	**+5.0**	71.8	61.4	53.6	**+10.3**	**+18.2**
gpt-6.1-terra	gpt-6-luna	29.7	4.3	53.6	+7.7	**+7.3**	60.9	66.7	73.0	−5.8	**−12.1**
gpt-6-luna	gpt-4.1	45.7	6.7	64.7	+7.7	**+6.3**	71.0	61.3	53.6	**+9.7**	**+17.4**
Claude Haiku	Claude Haiku	61.7	11.3	66.8	+9.8	**+5.7**	72.4	61.3	66.8	**+11.1**	**+5.7**
Claude Sonnet	Claude Haiku	65.0	9.7	66.8	+11.7	**+8.2**	75.0	60.4	75.7	**+14.6**	−0.7
gpt-6.1-terra	Qwen2.5-7B	31.0	4.0	50.3	+11.7	**+8.0**	58.3	62.3	73.0	−4.0	**−14.7**
gpt-6.1-terra	gpt-4o-mini	31.3	4.0	51.1	+12.8	**+9.6**	60.7	60.6	73.0	+0.1	**−12.3**
gpt-6-luna	gpt-6-luna	45.7	6.7	53.6	+13.8	**+11.1**	64.7	56.4	53.6	**+8.2**	**+11.1**
gpt-4o-mini	gpt-6-luna	67.0	30.3	53.6	+14.7	**+8.6**	62.1	56.1	51.1	+6.0	**+11.0**
gpt-4.1	gpt-6.1-terra	79.0	8.0	73.0	+15.7	**+7.3**	80.3	69.7	64.7	**+10.7**	**+15.7**
gpt-4o-mini	Qwen2.5-7B	67.0	30.3	50.3	+17.3	**+7.3**	57.7	53.3	51.1	+4.3	+6.6
gpt-4o-mini	gpt-4o-mini	67.0	30.3	51.1	+17.6	**+8.4**	59.6	52.4	51.1	+7.1	**+8.4**
gpt-6-luna	gpt-4o-mini	45.7	6.7	51.1	+19.1	**+13.7**	64.8	56.7	53.6	**+8.1**	**+11.2**
Claude Haiku	gpt-6-luna	61.7	11.3	53.6	+19.1	**+14.3**	67.9	63.0	66.8	+4.9	+1.1
gpt-6-luna	Qwen2.5-7B	48.0	6.7	50.3	+19.7	**+12.0**	62.3	58.0	53.6	+4.3	**+8.8**
gpt-4.1	Claude Haiku	79.0	8.0	66.8	+20.2	**+11.0**	77.8	63.7	64.7	**+14.1**	**+13.1**
gpt-4.1	gpt-4.1	79.0	8.0	64.7	+20.7	**+11.7**	76.3	63.3	64.7	**+13.0**	**+11.7**
Claude Haiku	gpt-4o-mini	61.7	11.3	51.1	+21.3	**+13.0**	64.1	57.9	66.8	+6.2	−2.7
Claude Sonnet	gpt-6-luna	65.0	9.7	53.6	+22.8	**+18.1**	71.7	67.8	75.7	+3.9	−4.0
Claude Sonnet	gpt-4o-mini	65.0	9.7	51.1	+28.1	**+15.9**	67.0	61.2	75.7	+5.8	**−8.7**
gpt-4.1	gpt-6-luna	79.0	8.0	53.6	+30.1	**+21.3**	74.9	61.2	64.7	**+13.7**	**+10.2**
gpt-4.1	gpt-4o-mini	79.0	8.0	51.1	+33.8	**+18.6**	69.7	57.9	64.7	**+11.8**	+5.0
gpt-4.1	Qwen2.5-7B	79.3	7.7	50.3	+36.7	**+16.0**	66.3	57.7	64.7	+8.7	+1.7

Table 28: Every detector–generator cell of H24 and H27 (ConflictBank, one-answer prompt, 300 conflict questions; %), sorted by the exploratory score. *Right*: share of questions on which the detector flags the true contradiction and names exactly its misinformation passage (the one definition of direction accuracy used throughout, computed on the detections that cell's generator saw); *Wrong*: share on which it marks a correct passage as suspect; *Gen. acc.*: the generator's accuracy with no flags; *Score*: right direction where the generator errs minus a correct passage marked where it is right; *Gain*: the LLM directional flag's change in points. *Hint*: the generator given the detector model's own plain answer, with no passage marked (H27); *det. alone*: the detector model's own plain accuracy. A detector's statistics can differ slightly between generators because concurrent runs made independent detector calls (up to 29 of 400 questions differ for one detector). Bold: significant under the paper's scheme (Gain: exploratory family; hint and det. alone columns: H27, Holm across cells).

correctly directed pair is enough to remove a misinformation document, so the detector removes the misinformation on 61.2% of conflict questions despite its low pair recall. We chose the document-level axes after seeing the LLM detector on the pilot split. The grid is a useful first estimate of a detector's effect, but with three detectors we cannot claim that one summary predicts better than the other.

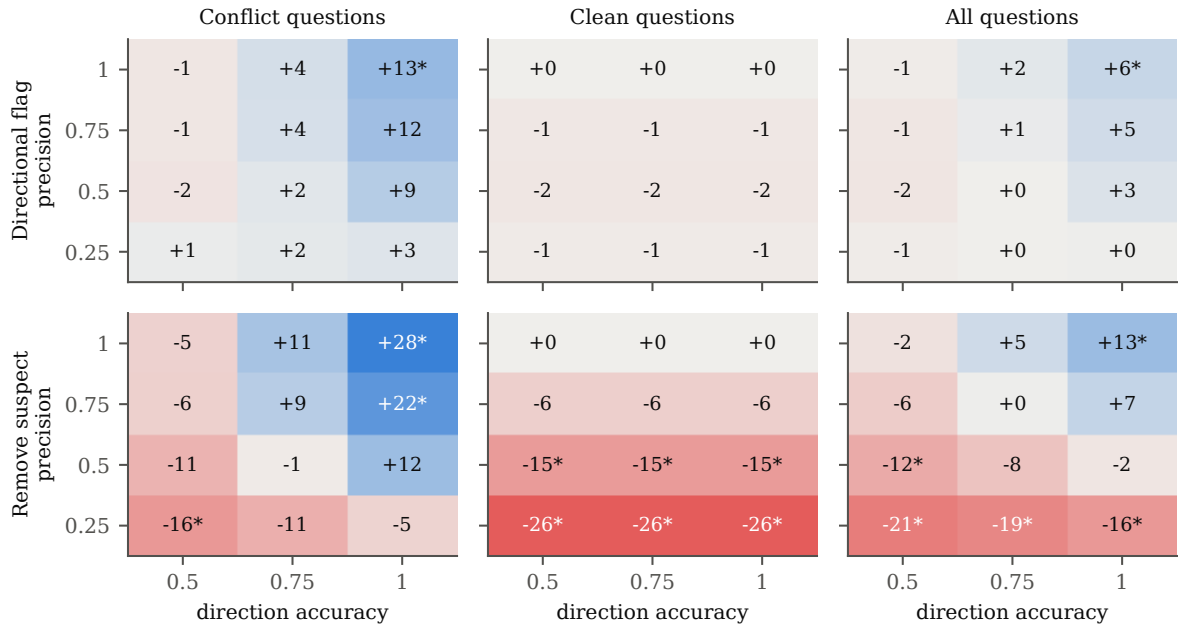


Figure 5: Synthetic detectors at recall 1: change in one-way strict accuracy (points) against no flags, by pair precision (rows) and direction accuracy (columns), for the directional flag (top) and removal (bottom). *: significant after Holm correction over the 36 cells of that mode and subset. Clean questions only see false flags, so their rows do not vary with direction accuracy.

Cell level (30 cells): accuracy with the flag, A	
$A \sim D + G$: coefficient of D	0.09 [0.03, 0.16]
coefficient of G	0.46 [0.36, 0.55]
$A \sim D$ + generator FE: coefficient of D	0.09 [0.02, 0.15]
exact permutation of D across detectors	$p = 0.29$
Leave one detector out: R^2 for the gain	
G only	0.17
$D + G$	-0.07
gap $D - G$ (registered)	-0.21
mechanism score	0.62
Question level (9000 cell-questions), points	
detector right on q	$+34.5$ [$+32.8$, $+36.1$]
detector marks a correct passage	-31.6 [-33.9, -29.4]
generator right on its own	$+37.9$ [$+36.3$, $+39.5$]
Calibration: gain $\sim$ score	
slope	0.64 [0.54, 0.74]

Generator	Follows a needed right flag	Follows a wrong-side flag
Qwen2.5-7B	61 (327)	77 (71)
Claude Haiku	57 (189)	83 (71)
gpt-4.1	69 (171)	83 (60)
gpt-4o-mini	66 (446)	76 (75)
gpt-6-luna	75 (281)	84 (58)
gpt-6.1-terra	68 (154)	88 (97)

Table 29: H24 without the shared $-G$ term (§15.4, re-analysis of existing cells). D: the detector's direction accuracy; G: the generator's no-flag accuracy. CIs: 2,000 question resamples, shared across cells. The question-level model is a linear probability model with detector and generator fixed effects and a random intercept per question. Lower table: share of questions the generator gets right when the detector is right and the generator alone is wrong in every seed, and share it gets wrong when the detector marks a correct passage and the generator alone is right in every seed (count of cell-questions in brackets).

Setting	Neutral flag	Setting	Neutral flag
CB single, gpt-4o-mini	−18.7	RAMDocs, gpt-6.1-terra	−0.7
CB minimal, gpt-4o-mini	−18.1	CB forced, gpt-6.1-terra	−0.3
RAMDocs, gpt-4.1 (minimal)	−11.0	CB forced, gpt-4o-mini	−0.2
CB single, gpt-6-luna	−9.4	RAMDocs, Claude Sonnet	0.0
CB multi, gpt-4o-mini	−8.8	RAGuard, gpt-6-luna	0.0
RAMDocs, gpt-4.1 (default)	−8.6	RAGuard, Claude Haiku	+0.3
CB multi, Qwen2.5-7B	−8.3	CB single, Claude Haiku	+0.7
CB single, gpt-4.1	−8.0	CB majority, gpt-4o-mini	+1.1
RAMDocs, gpt-4o-mini (default)	−6.7	RAMDocs, Claude Haiku	+1.2
CB multi, gpt-4.1	−6.3	CB forced, gpt-6-luna	+1.6
RAMDocs, Qwen2.5-7B (default)	−5.8	RAGuard, gpt-6.1-terra	+2.4
RAMDocs, Qwen2.5-7B (minimal)	−5.8	CB single, Claude Sonnet	+3.0
RAMDocs, gpt-4o-mini (minimal)	−5.3	RAGuard, gpt-4o-mini	+3.2
CB majority, gpt-6.1-terra	−3.7	CB forced, Claude Haiku	+3.9
CB single, Qwen2.5-7B	−3.3	CB majority, Claude Haiku	+4.3
RAMDocs, gpt-6-luna	−3.1	RAGuard, Claude Sonnet	+4.4
CB single, gpt-6.1-terra	−2.3	CB majority, gpt-6-luna	+6.6
Meta-regression (points; CIs from the question-cluster bootstrap)			
intercept (OpenAI, unknown allowed, synthetic)	−5.7 [−7.0, −4.4]		
Anthropic	+6.3 [+4.7, +7.7]		
Qwen	−1.2 [−3.9, +1.4]		
forced (no unknown)	+5.4 [+3.5, +7.1]		
natural data	+2.1 [+0.4, +3.7]		

Table 30: The 34 settings pooled for the oracle neutral flag (conflict questions, change in points). The pooled estimate was registered as exploratory in §13.4. Pooled: −2.8 (DerSimonian–Laird); 95% CI from a bootstrap that resamples questions within each benchmark jointly across the settings that share them [−3.5, −1.9]; prediction interval for a new setting [−12.9, +7.4]. Moderators (§15.4): model family (reference OpenAI), no "unknown" allowed (forced prompt) and natural (RAMDocs, RAGuard) rather than generated data (ConflictBank).

	Neutral flag, utility change, cost c of misinformation				Dir. flag	Neutral flag: plain answer → flagged answer		
Setting	$c=0$	$c=1$	$c=2$	$c=4$	$c=2$	misinfo→abstain	misinfo→correct	correct→abstain
CB single, gpt-4o-mini	**−18.7**	−5.3	+8.0	**+34.7**	**+83.7**	45	7	37
CB single, gpt-6-luna	**−9.4**	−2.6	+4.3	+18.1	**+85.1**	27	12	20
CB single, gpt-6.1-terra	−2.3	0.0	+2.3	+7.0	**+51.7**	2	26	4
CB single, Claude Haiku	+0.7	+0.6	+0.4	+0.2	**+66.9**	0	25	0
CB single, Claude Sonnet	+3.0	+6.7	+10.3	+17.7	**+45.7**	2	34	0
CB forced, gpt-4o-mini	−0.2	−0.6	−0.9	−1.6	**+74.0**	1	19	0
CB majority, gpt-4o-mini	+1.1	**+24.6**	**+48.0**	**+94.9**	**+118.2**	26	11	22
RAGuard, gpt-4o-mini	+3.2	+6.6	+10.0	+16.8	**+34.6**	12	7	8
RAGuard, gpt-6-luna	0.0	−1.1	−2.3	−4.5	**+29.3**	10	6	8
RAGuard, Claude Haiku	+0.3	−1.8	−4.0	−8.3	**+26.5**	7	4	11
RAMDocs, gpt-6-luna	**−3.1**	−7.7	−12.2	−21.3	**+160.9**	0	1	0
RAMDocs, Claude Haiku	+1.2	+9.1	+17.0	**+32.9**	**+125.2**	0	9	0

Table 31: Abstention-aware view of the oracle flags (conflict questions; §15.4, re-analysis). Utility per answer: +1 correct, $-c$ if it contains the misinformation, 0 for abstaining or any other wrong answer; $c = 0$ is strict accuracy. Bold: significant in one exploratory Holm family per c over all 28 settings. Right: of the question–seed pairs whose plain answer contained the misinformation (or was correct), the share whose flagged answer abstained or was correct.